

Mid term Exam for Financial Econometrics with Python

PRAT Paul; GAVINI Charles; FOURNIER Justin; BLANC Mathieu

November 13, 2024

Contents

1	Introduction	1
2	Preliminary	1
2.1	AMAZON	1
2.2	Data Table	1
2.3	Checking the 25 Years range condition	2
3	First Results	2
3.1	Prices Evolutions	2
3.2	Calculating Returns	3
3.3	Squared Returns	3
4	Amazon and the 8 Stylized Facts	4
4.0.1	Summary statistics	4
4.1	Prices are non-stationary	4
4.2	Returns are stationary	5
4.3	Asymmetry	5
4.4	Heavy tails	6
4.5	Gaussianity	7
4.5.1	High frequency non-Gaussianity	7
4.5.2	Aggregational Gaussianity	7
4.6	Returns are not autocorrelated	7
4.7	Volatility clustering and long range dependence of squared returns	7
4.8	Leverage effect	7
A	Appendix: Python Code	7
B	To go further, CAPM pricing model	8

1 Introduction

This document provides a comprehensive presentation of our results, including all relevant tables, figures, and calculations. The report is structured into distinct parts, beginning with the importation of essential Python libraries. We then initialize variables to organize the data into different categories (e.g., daily, monthly, returns, log returns), allowing for clear analysis and comparison across various data types and intervals.

2 Preliminary

2.1 AMAZON

The selected stock for this analysis is Amazon due to its significant relevance in current global markets, its impressive growth over time and its position as a major industry leader. The ticker from yahoo finance is "**AMZN**" on the Nasdaq stock exchange [AMAZON on Yahoo Finance](#) First, importing the Amazon stock with yfinance, then display the pandas table. We will import 25 years, 8 months and 25 days of data (from 1999-01-21 to 2024-10-16).

2.2 Data Table

The data printed here is the preview of the Amazon stock extraction from yahoo finance:

Date	Open	High	Low	Close	Adj Close	Volume
1999-01-21	2.612500	2.759375	2.314063	2.650000	2.650000	940964000
1999-01-22	2.487500	3.146875	2.468750	3.075000	3.075000	875316000
1999-01-25	3.037500	3.084375	2.750000	2.809375	2.809375	546476000
1999-01-26	2.815625	3.031250	2.765625	2.877344	2.877344	490696000
1999-01-27	3.353125	3.493750	3.000000	3.140625	3.140625	700452000

Table 1: Preview of Amazon Stock Data from Yahoo Finance

2.3 Checking the 25 Years range condition

We need to verify that the data displays accurately over the 25 years range. Fortunately, the extracted Amazon data has been available since January 1999. To ensure the data's continuity and completeness, we will implement a Python script that identifies and counts any gaps within the dataset. By visualizing the dates of these gaps, we can easily detect any significant interruptions that could potentially impact our data analysis

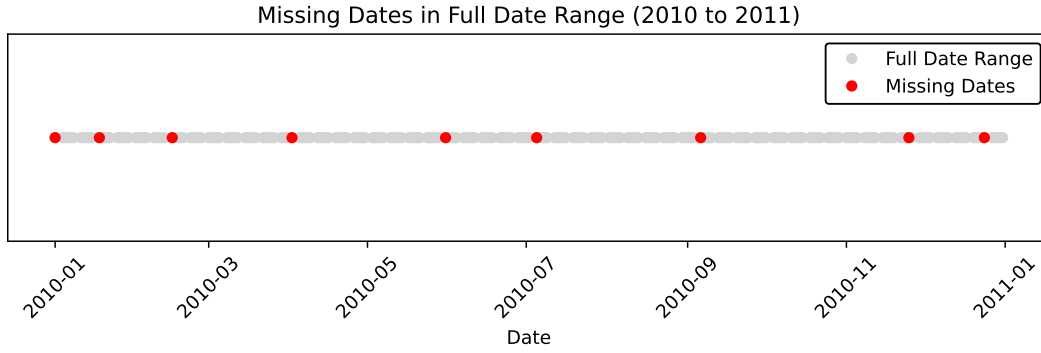


Figure 1: Missing Dates in Full Date Range (2010 to 2011)

We identified a total of 238 isolated days of data gaps per year across the 25-years range (6476 values). Therefore, the data remains reliable for our stylized facts analysis. The missing data points in our dataset are randomly distributed and account for 3.7% of the total data. According to scientific studies on data reliability for volatility testing, a dataset with up to 10 [2]

3 First Results

3.1 Prices Evolutions

With the accuracy and the reliability of our dataset confirmed, we begin by plotting the evolution of prices over 4 different periods : Daily, Weekly, Monthly and Yearly prices.




Figure 2: Prices over time by frequency

3.2 Calculating Returns

Using the processed data, we can now output graphs for several key metrics: daily prices, daily log prices, daily simple returns, and daily log returns. Plotting these metrics will allow us to observe daily price movements, the transformation of prices into *log* form for trend analysis, as well as daily returns and their logarithmic equivalents.

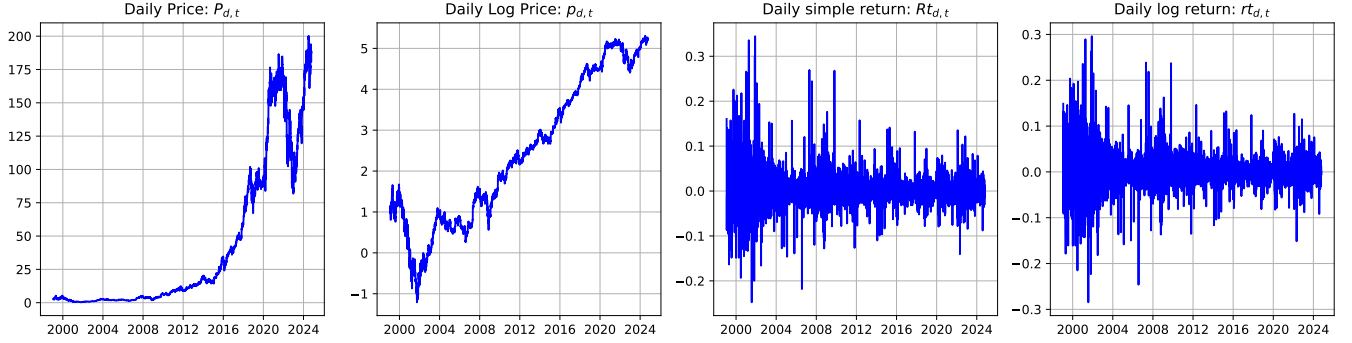


Figure 3: Prices, returns and log returns

3.3 Squared Returns

To complete the analysis of daily price data, we also plot the daily squared returns and daily squared log returns, (providing us a key insight on the volatile behavior on the potential stock risk of our data set.)

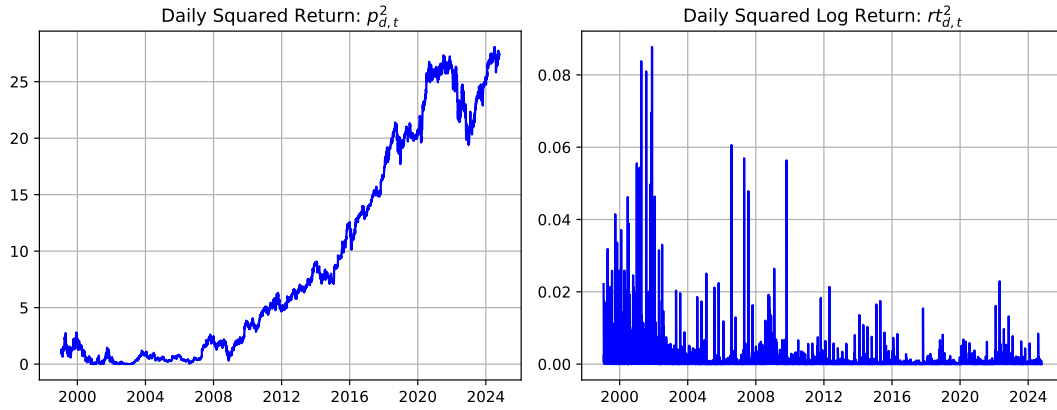


Figure 4: Squared daily returns and daily log returns

4 Amazon and the 8 Stylized Facts

4.0.1 Summary statistics

	daily	weekly	monthly	annual
Mean	0.06351	0.31014	1.32164	22.85692
St.Deviation	3.24131	6.76830	13.06275	45.28052
Diameter.C.I.Mean	0.07896	0.36213	1.45887	18.11598
Skewness	0.41992	0.05008	-0.45895	-0.15137
Kurtosis	11.15158	7.60655	2.59462	-0.64793
Excess.Kurtosis	8.15158	4.60655	-0.40538	-3.64793
Min	-28.45678	-38.51804	-53.02674	-68.54809
Quant5	-4.61051	-9.74288	-20.16713	-55.72274
Quant25	-1.25994	-2.64062	-4.98163	-7.23999
Median	0.04108	0.30519	2.09626	23.07665
Quant75	1.39659	3.40897	8.45973	55.96192
Quant95	4.47118	10.67416	20.90661	94.77653
Max	29.61811	56.11507	48.35221	102.44636
Jarque.Bera.stat	33735.75720	3235.87866	97.20696	0.51147
Jarque.Bera.pvalue.X100	0.00000	0.00000	0.00000	77.43487
Lillie.test.stat	0.10194	0.09591	0.08194	0.06494
Lillie.test.pvalue.X100	0.10000	0.10000	0.10000	99.00000
N.obs	6474.00000	1342.00000	308.00000	24.00000

Table 2: Summary statistics for the amazon stock

4.1 Prices are non-stationary

The first feature that will highlight non-stationarity of the prices is the comparison of p_t vs p_{t-1} .

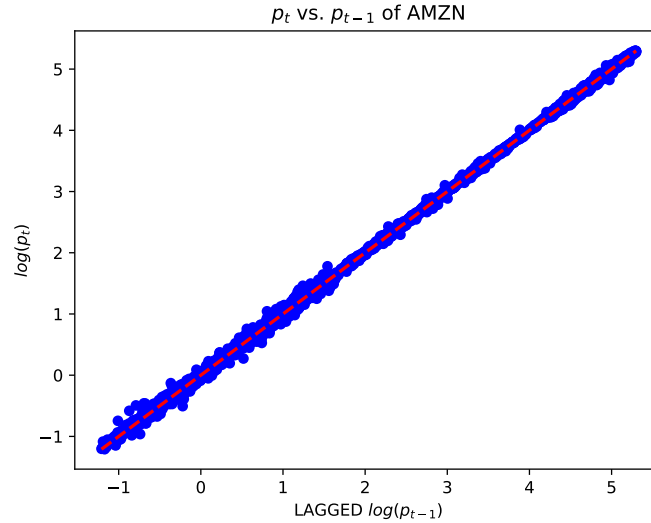


Figure 5: Comparison of $\log(p_t)$ vs $\log(p_{t-1})$

The graph in Figure 5 demonstrates this strong linear relationship, indicating that Amazon's prices at time t are highly dependent on those at $t - 1$ and lack mean reversion, supporting the idea of non-stationarity. Additionally, the empirical autocorrelation function (ACF) of Amazon's daily prices shows a slow decay, further suggesting non-stationarity, as shown in the next figure.




Figure 6: Autocorrelations of daily and monthly Prices (1999-2024)

4.2 Returns are stationary




Figure 7: Daily log-returns (1999 - 2024)

The graph Figure 7 illustrates the daily log-returns of Amazon's stock over a 25-year period, from 1999 to 2024. *log*-returns are a common way to measure the percentage change in stock prices, and they help assess the stability or stationarity of the returns over time. In a stationary series, we would expect the properties, such as mean and variance, to remain constant over time. However, here we observe significant differences in volatility across the timeline.

In the early years (around 1999-2005), there is noticeably higher volatility in Amazon's log-returns, with frequent large spikes both upwards and downwards. This period corresponds to the tech boom and subsequent dot-com bubble burst, during which many tech stocks, including Amazon, experienced extreme price fluctuations. Additionally, as a relatively new and fast-growing company, Amazon's stock likely faced higher uncertainty and speculative trading, contributing to greater volatility.

4.3 Asymmetry

TODO Stylized Fact 3: The distribution of returns is asymmetric and often negatively skewed, reflecting the fact that the downturns of financial markets are often much steeper than the recoveries. Investors tend to react more strongly to negative news than to positive news.

	daily	weekly	monthly	annual
Skewness	0.41992	0.05008	-0.45895	-0.15137
Kurtosis	11.15158	7.60655	2.59462	-0.64793

Table 3: Skewness and kurtosis for *log* returns




Figure 8: Rolling skewness

4.4 Heavy tails

As showcased in the Table 3, there is a large excess kurtosis

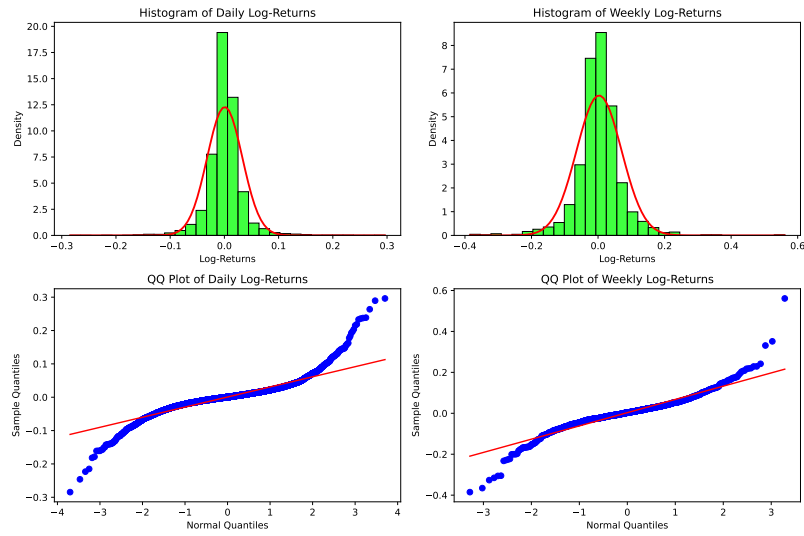


Figure 9: Histogram and QQ plot of log-returns

Here, the QQ-Plots explicit clearly how our sample distinguishes from the normal distribution. The QQ-plot provide graphical evidence that the tails of the daily returns distributions are heavier than the tails of the normal distribution as: The points on the left of the graph which represent the lower quantiles (i.e. the points in the left tail of the empirical distribution) are below the blue line. The lower quantiles of the empirical distribution are much smaller than what you should expect from a Normal random variable with the same empirical mean and standard deviation of the sample the left tail of the empirical distribution is heavier than one of a Normal Distribution. similar conclusions for the right tail.




Figure 10: QQ-plots for $\nu = 10, 5, 3$

4.5 Gaussianity

TODO

4.5.1 High frequency non-Gaussianity

As the result in Table 2 the skewness is positive for daily and monthly data [1] TODO add an article about amzn skewness. The 3rd central moment is defined as $\mu_3 := E((X - m_1)^3)$. The skewness of r_t is defined as:

$$\text{Skew}(r_t) := E \left[\left(\frac{X - m_1}{\sigma} \right)^3 \right] = \frac{\mu_3}{\sigma^3} = \frac{\mu_3}{\mu_2^{3/2}}.$$

In this case, and as emphasized in the 3rd stylized fact, $\text{Skew}(r_t) > 0$, large realizations of X are more often larger than the mean μ . Skewness is thus used as a measure of asymmetry of the distribution $f_X(x)$. Therefore:

$\text{Skew}(r_t) > 0$, so the distribution is said to be **right skewed**.

$\text{Skew}(r_t) > 0$, then $\mu > \text{median}$.

4.5.2 Aggregational Gaussianity




Figure 11: Log returns $r_t := p_t - p_{t-1}$: **annual** “adjusted closing” of S&P500. Sample: 01-21-1999 to 10-16-2024. **Left panel:** empirical and Normal cdf’s for the standardized annual returns of the S&P500. **Right panel:** values of $|G_T(r_t) - \Phi(r_t, \hat{\mu}, \hat{\sigma}^2)|$ (blue line) and critical values for the Lilliefors test for the three significance levels 10%, 5% and 1%.

The blue line is under the critical values lines, So the test is respected and so for the Gaussianity.

4.6 Returns are not autocorrelated

TODO

4.7 Volatility clustering and long range dependence of squared returns

TODO

4.8 Leverage effect

TODO

A Appendix: Python Code

Below is the Python code used in this analysis.

```

1  # Python code example
2  import numpy as np
3  import pandas as pd
4
5  def analyze_data(data):
6      mean = np.mean(data)
7      std_dev = np.std(data)
8      return mean, std_dev
9
10 data = [1, 2, 3, 4, 5]
11 mean, std_dev = analyze_data(data)
12 print(f"Mean: {mean}, Standard Deviation: {std_dev}")

```

Listing 1: Python Code for Analysis

B To go further, CAPM pricing model

References

- [1] John Doe and Jane Smith. An example article. *Journal of Examples*, 42(1):1–10, 2023.
- [2] Giovanni Pumi et al. Estimation of long-range dependent models with missing data: to impute or not to impute? *arXiv preprint*, 2023.