

Mid term Exam for Financial Econometrics with Python

PRAT Paul; GAVINI Charles; FOURNIER Justin; BLANC Mathieu

November 14, 2024

Contents

1	Introduction	1
2	Preliminary	1
2.1	AMAZON	1
2.2	Data Table	1
2.3	Checking the 25 Years range condition	1
3	First Results	2
3.1	Prices Evolutions	2
3.2	Calculating Returns	2
4	Amazon and the 8 Stylized Facts	3
4.0.1	Summary statistics	3
4.1	Prices are non-stationary	3
4.2	Returns are stationary	4
4.3	Asymmetry	5
4.4	Heavy tails	5
4.5	Gaussianity	6
4.5.1	High frequency non-Gaussianity	6
4.5.2	Aggregational Gaussianity	6
4.6	Returns are not autocorrelated	6
4.7	Volatility clustering and long range dependence of squared returns	7
4.8	Leverage effect	8
A	References	9
B	Python code	9

1 Introduction

This document provides a comprehensive presentation of our results, including all relevant tables, figures, and calculations. The report is structured into distinct parts, beginning with the importation of essential Python libraries. We then initialize variables to organize the data into different categories (e.g., daily, monthly, returns, log returns), allowing for clear analysis and comparison across various data types and intervals.

2 Preliminary

2.1 AMAZON

The selected stock for this analysis is Amazon due to its significant relevance in current global markets, its impressive growth over time and its position as a major industry leader. The ticker from yahoo finance is "**AMZN**" on the Nasdaq stock exchange [AMAZON on Yahoo Finance](#) First, importing the Amazon stock with yfinance, then display the pandas table. We will import 25 years, 8 months and 25 days of data (from 1999-01-21 to 2024-10-16).

2.2 Data Table

The data printed here is the preview of the Amazon stock extraction from yahoo finance:

Date	Open	High	Low	Close	Adj Close	Volume
1998-12-30	2.775000	2.860417	2.532292	2.677083	2.677083	651672000
1998-12-31	2.643750	2.758333	2.634375	2.677083	2.677083	365964000
1999-01-04	2.730729	2.966667	2.665625	2.957813	2.957813	785844000
1999-01-05	2.739063	3.243750	2.662500	3.112500	3.112500	1257464000
1999-01-06	3.409375	3.509375	3.350000	3.450000	3.450000	723532000

Table 1: Preview of Amazon Stock Data (5 first datas) from "AMZN" in Yahoo Finance

2.3 Checking the 25 Years range condition

We need to verify that the data displays accurately over the 25 years range. Fortunately, the extracted Amazon data has been available since January 1999. To ensure the data's continuity and completeness, we will implement a Python script that identifies and counts any gaps within the dataset. By visualizing the dates of these gaps, we can easily detect any significant interruptions that could potentially impact our data analysis




Figure 1: Missing Dates in a partial date range (01-01-2010 to 01-01-2011)

We identified a total of 238 isolated days of data gaps per year across the 25-years range (6476 values). Therefore, the data remains reliable for our stylized facts analysis. The missing data points in our dataset are randomly distributed and account for 3.7% of the total data. According to scientific studies on data reliability for volatility testing, a dataset with up to 10% missing data is considered reliable for statistical testing. [2]

3 First Results

3.1 Prices Evolutions

With the accuracy and the reliability of our dataset confirmed, we begin by plotting the evolution of prices over 4 different frequency : Daily, Weekly, Monthly and Yearly prices.

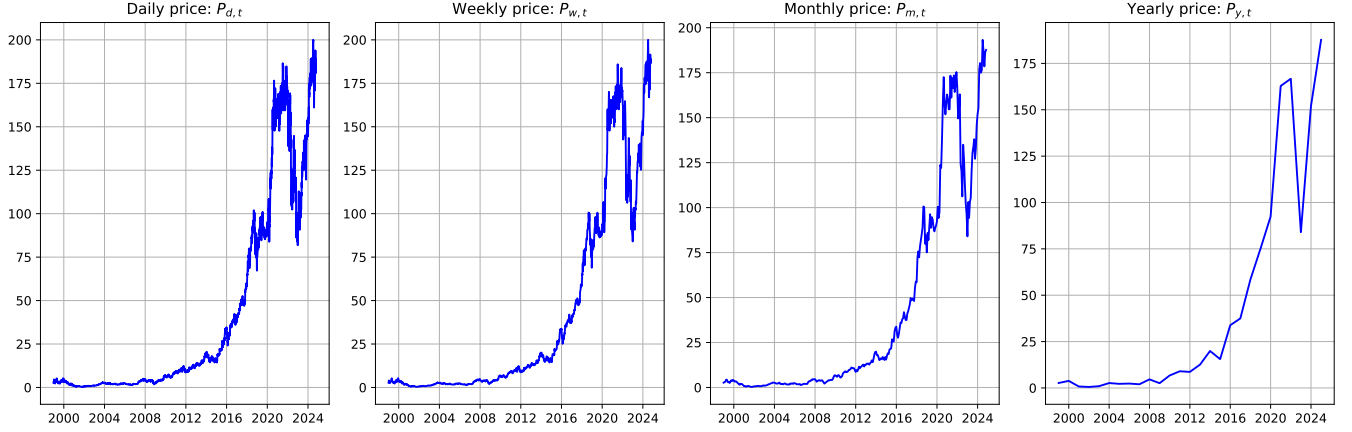


Figure 2: Prices over time P_t by frequency daily, weekly, monthly and annual the AMZN stock. Sample: **01-21-1999** to **10-16-2024**.

3.2 Calculating Returns

Using the processed data, we can now output graphs for several key metrics: daily prices, daily log prices, daily simple returns, and daily log returns. Plotting these metrics will allow us to observe daily price movements, the transformation of prices into *log* form for trend analysis, as well as daily returns and their logarithmic equivalents.

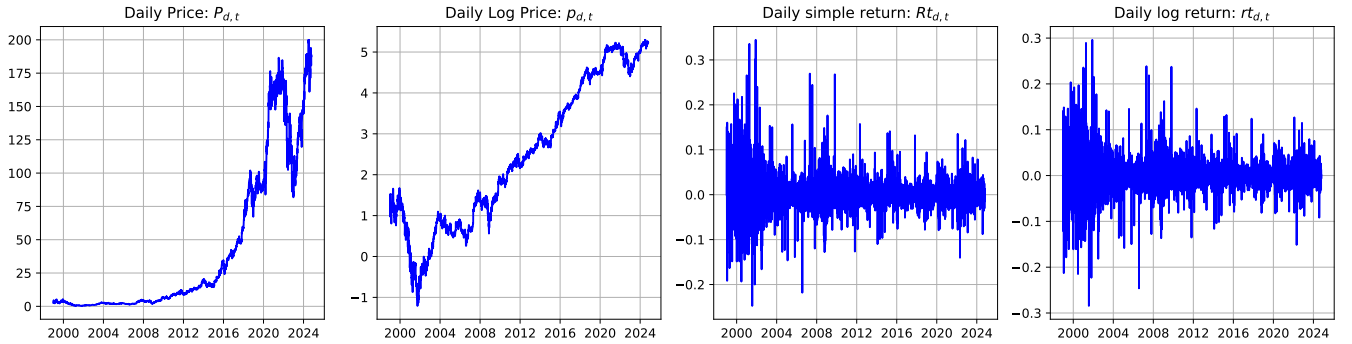


Figure 3: Prices P_t , returns R_t and log returns r_t of the AMZN stock. Sample: **01-21-1999** to **10-16-2024**.

4 Amazon and the 8 Stylized Facts

4.0.1 Summary statistics

	daily	weekly	monthly	annual
Mean	0.06551	0.28602	1.34694	15.59259
St.Deviation	3.27670	6.78240	13.04915	56.87533
Diameter.C.I.Mean	0.07973	0.36248	1.45499	22.29513
Skewness	0.39404	0.04813	-0.46401	-0.99033
Kurtosis	11.04424	7.51963	2.60379	1.46124
Excess.Kurtosis	8.04424	4.51963	-0.39621	-1.53876
Min	-28.45678	-38.51804	-53.02674	-158.75126
Quant5	-4.64852	-9.91694	-20.10778	-66.66688
Quant25	-1.26082	-2.66403	-4.96881	-17.23374
Median	0.04153	0.30015	2.12289	21.13060
Quant75	1.39904	3.40521	8.48954	55.72364
Quant95	4.50385	10.67205	20.87667	94.22164
Max	29.61811	56.11507	48.35221	102.44636
Jarque.Bera.stat	33141.87111	3169.38203	98.37746	6.31063
Jarque.Bera.pvalue.X100	0.00000	0.00000	0.00000	4.26249
Lillie.test.stat	0.10370	0.09576	0.08174	0.09522
Lillie.test.pvalue.X100	0.10000	0.10000	0.10000	79.74330
N.obs	6488.00000	1345.00000	309.00000	25.00000

Table 2: Summary statistics for the AMZN stock. Sample: **01-21-1999** to **10-16-2024**.

4.1 Prices are non-stationary

The first feature that will highlight non-stationarity of the prices is the comparison of p_t vs p_{t-1} .

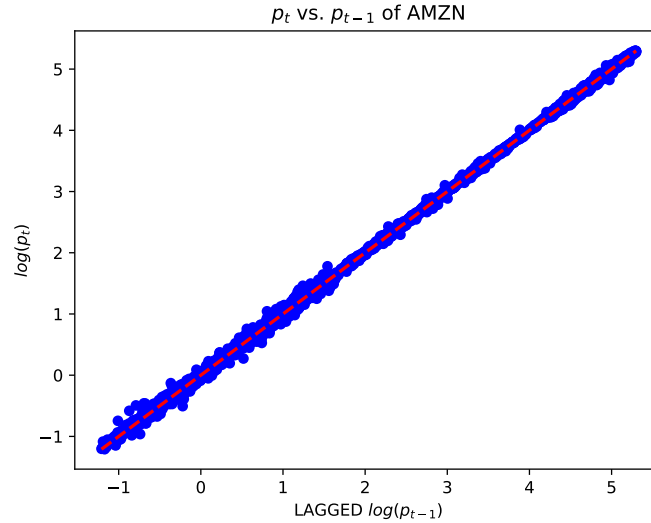


Figure 4: Comparison of $\log(p_t)$ vs $\log(p_{t-1})$ of the AMZN stock. Sample: **01-21-1999** to **10-16-2024**.

The graph in Figure 4 demonstrates this strong linear relationship, indicating that Amazon's prices at time t are highly dependent on those at $t - 1$ and lack mean reversion, supporting the idea of non-stationarity. Additionally, the empirical autocorrelation function (ACF) of Amazon's daily prices shows a slow decay, further suggesting non-stationarity, as shown in the next figure.

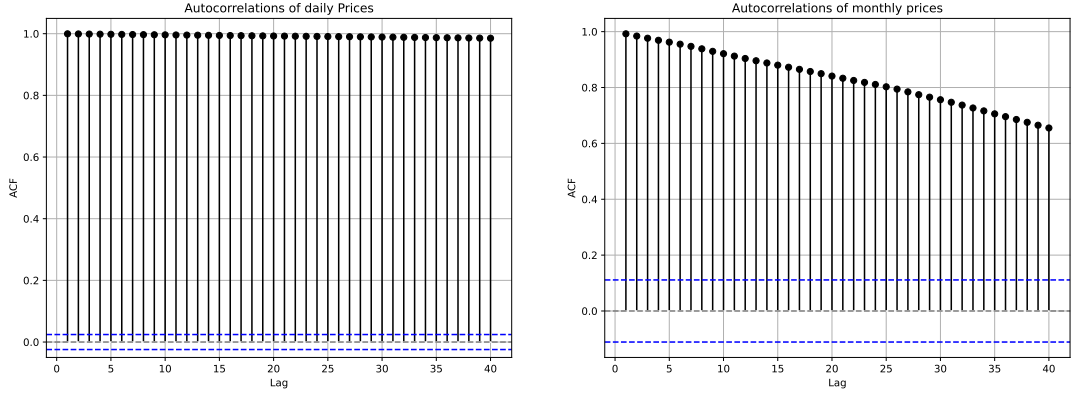


Figure 5: Autocorrelations of daily and monthly prices of the AMZN stock. Sample: **01-21-1999** to **10-16-2024**.

For the amzon daily and monthly prices time series, we expect to see large values of $\hat{\rho}_k$, near to 1, slowly decaying as k increases this is the **long memory property**.

4.2 Returns are stationary

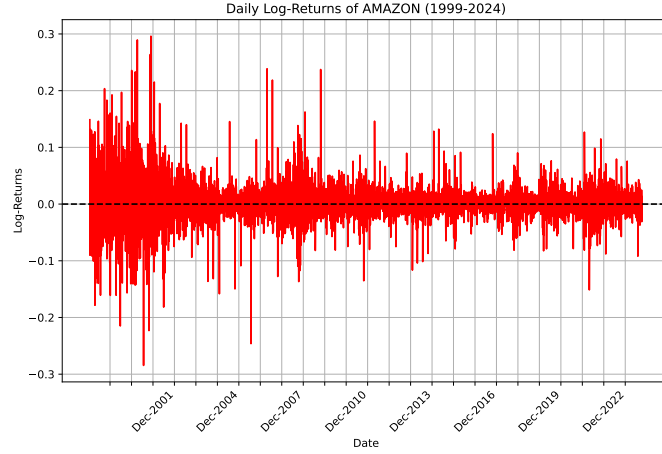


Figure 6: Daily Log-returns $r_t := p_t - p_{t-1}$ of the AMZN stock. Sample: **01-21-1999** to **10-16-2024**.

Log-returns are a common way to measure the percentage change in stock prices, and they help assess the stability or stationarity of the returns over time. In a stationary series, we would expect the properties, such as mean and variance, to remain constant over time. However, here we observe significant differences in volatility across the timeline.

In the early years (around 1999-2005), there is noticeably higher volatility in Amazon's log-returns, with frequent large spikes both upwards and downwards. This period corresponds to the tech boom and subsequent dot-com bubble burst, during which many tech stocks, including Amazon, experienced extreme price fluctuations. Additionally, as a relatively new and fast-growing company, Amazon's stock likely faced higher uncertainty and speculative trading, contributing to greater volatility.

4.3 Asymmetry

	daily	weekly	monthly	annual
Skewness	0.39404	0.04813	-0.46401	-0.99033
Kurtosis	11.04424	7.51963	2.60379	1.46124

Table 3: **Skewness and kurtosis of daily, weekly, monthly and annual \log returns of the AMZN stock. Sample: 01-21-1999 to 10-16-2024.**




Figure 7: **Rolling skewness** of AMZN stock. Sample: **01-21-1999 to 10-16-2024**. The **red bands** corresponds to the limit of acceptance, the blue line correspond to the rolling skewness with $T=252$

For this case, Table 3 highlights that the skewness of daily returns for AMZN stock is positive. This does not confirm stylized fact 3, as it is not very common; however, it indicates that the mean return is higher than the median of the sample [1]. Consequently, Amazon investors tend to experience steeper upturns than downturns, and they react more positively to good news than they react negatively to bad news. Looking at the rolling skewness of simple returns in Figure 7, we can clearly see that the skewness (calculated over a 252-day interval) varies significantly depending on the interval, and has been very negative at times (e.g., December 2004) but is generally positive.”

4.4 Heavy tails

As showcased in the Table 3, there is a large excess kurtosis

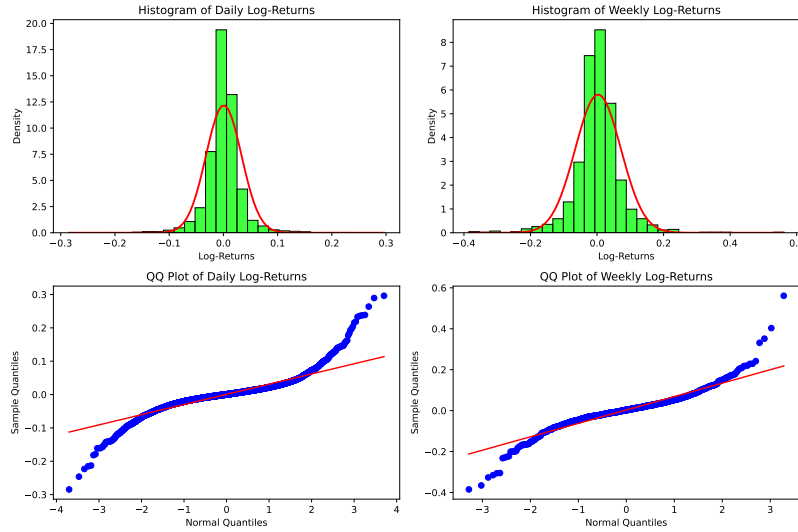


Figure 8: Log returns $r_t := p_t - p_{t-1}$: **histograms of daily, monthly “adjusted closing” of AMAZON. Sample: 01-21-1999 to 10-16-2024**. QQ plot against quantiles of normal distribution with same mean and variance as the empirical distribution of returns.

Here, the QQ-plots clearly illustrate how our sample differs from the normal distribution. The QQ-plot provides graphical evidence that the tails of the daily returns distribution are heavier than those of the normal distribution, as follows: the points on the left side of the graph, representing the lower quantiles (i.e., the left tail of the empirical distribution), fall below the blue line. This indicates that the lower quantiles of the empirical distribution are much smaller than expected for a normal random variable with the same empirical mean and standard deviation.




Figure 9: Log returns $r_t := p_t - p_{t-1}$: **daily** “adjusted closing” of AMZN stock. Sample: **01-21-1999 to 10-16-2024**. QQ plot of Sample standardized quantiles (0 mean and unit variance) of daily log-returns against quantiles of standardized (0 mean and unit variance) Student-t distributions with $\nu = 10, 5$, and 3 degrees of freedom.

4.5 Gaussianity

4.5.1 High frequency non-Gaussianity

The aggregate gaussianity, states that lower frequency returns (monthly) tend to be Gaussian (symmetric about the mean) even if higher frequency returns (daily) are not. To test this stylized fact we perform a Jarge-Bera test. The null hypothesis and its alternative are in the Table 2.

4.5.2 Aggregational Gaussianity




Figure 10: Log returns $r_t := p_t - p_{t-1}$: **annual** “adjusted closing” of AMZN. Sample: **01-21-1999 to 10-16-2024**. **Left panel**: empirical and Normal cdf’s for the standardized annual returns of AMZN. **Right panel**: values of $|G_T(\tilde{r}_t) - \Phi(\tilde{r}_t, \hat{\mu}, \hat{\sigma}^2)|$ (blue line) and critical values for the Lilliefors test for the three significance levels 10%, 5% and 1%.

The blue line is under the critical values lines, So the test is respected and so for the Gaussianity.

4.6 Returns are not autocorrelated

Stylized fact 6 posits that returns are not autocorrelated. Autocorrelation in a weakly stationary process measures the correlation between values of the process at different time points.




Figure 11: **Empirical Autocorrelation (ACF)** of r_t , the daily, weekly and monthly log-return from the “adjusted closing” of AMAZON from Yahoo Finance. the autocorrelation of order 0 is not reported to have a better graphical representation of the smaller autocorrelations of order k . Sample: **01-21-1999 to 10-16-2024**.

We will now test if the autocorrelations are 0, i.e. they are insignificant, through 2 Q-tests: the Box-Pierce test and the Ljung-Box test.

	lag	acf	acf diam.	acf test	B-P stat	B-P pval	L-B stat	L-B pval	crit
0	1	-0.009000	0.024000	-0.719000	0.517000	0.472000	0.517000	0.472000	3.841000
4	5	-0.016000	0.024000	-1.290000	13.110000	0.022000	13.119000	0.022000	11.070000
14	15	-0.022000	0.024000	-1.803000	26.089000	0.037000	26.124000	0.037000	24.996000
24	25	0.016000	0.024000	1.290000	62.020000	0.000000	62.177000	0.000000	37.652000

Table 4: Empirical Autocorrelation (ACF), ACF “diameter” ($1.96 * \sqrt{\frac{1}{T}}$), Box-Pierce test and Ljung-Box test:statistics and p-values. Data : the daily log-returns of Amazon stock from **01-21-1999 to 10-16-2024** (source : **Yahoo Finance**)

	lag	acf	acf diam.	acf test	B-P stat	B-P pval	L-B stat	L-B pval	crit
0	1	0.054000	0.111000	0.950000	0.903000	0.342000	0.911000	0.340000	3.841000
4	5	0.022000	0.111000	0.386000	8.178000	0.147000	8.329000	0.139000	11.070000
14	15	-0.052000	0.111000	-0.917000	23.761000	0.069000	24.481000	0.057000	24.996000
24	25	-0.056000	0.111000	-0.983000	38.351000	0.043000	40.250000	0.027000	37.652000

Table 5: Empirical Autocorrelation (ACF), ACF “diameter” ($1.96 * \sqrt{\frac{1}{T}}$), Box-Pierce test and Ljung-Box test:statistics and p-values. Data : the monthly log-returns of Amazon stock from **01-21-1999 to 10-16-2024** (source : **Yahoo Finance**)

Daily returns : ACF are significantly different from 0 ($ACF \geq ACF \text{ diameter}$) Monthly returns : ACF are NOT significantly different from 0 ($ACF \leq ACF \text{ diameter}$)

For daily returns, the BP test shows very small but statistically significant autocorrelations, indicating weak dependencies that are practically negligible.

For monthly returns, the BP test yields high p-values, confirming no significant autocorrelation, which aligns with the idea that returns are not autocorrelated over time.

4.7 Volatility clustering and long range dependence of squared returns

Volatility clustering is a phenomenon where periods of high market volatility are often followed by high volatility, and vice versa. To capture and analyze this phenomenon, financial models such as ARCH (Autoregressive Conditional Heteroskedasticity) and GARCH are commonly used. We can easily perceive it on the graph below, (from december 2001 to december 2004) phase of low volatility.




Figure 12: Rolling standard deviation from the “adjusted closing” of AMZN. Sample: **01-21-1999** to **10-16-2024**.

This persistence in the autocorrelation of squared returns reflects volatility clustering. High volatility often persists over time before settling into a lower volatility regime; this is how time dependance is reflected.




Figure 13: Autocorrelations of the daily, weekly and monthly absolute log-returns r_t from the “adjusted closing” of AMZN. Sample: **01-21-1999** to **10-16-2024**. The blue dotted bands represents the confidence intervals (Barlett intervals), $\frac{1}{\sqrt{T}}$ where T is the number of samples.

We observe that the autocorrelation is continuous, as indicated by the trendline, which aligns with the previous graph. Additionally, it becomes apparent that as the time interval changes (from daily to weekly, monthly, and annually) the autocorrelation becomes more pronounced between intervals. This aligns with the volatility clustering phenomenon discussed earlier. This effect occurs because ARCH and GARCH models are sensitive to sampling frequency, with their impact being more noticeable at shorter frequencies (daily, weekly) than at longer ones (monthly, annually).

4.8 Leverage effect

The leverage effect demonstrates the negative correlation between an asset’s returns and its volatility. Figure 14 shows this effect through cross-correlation values over time. The graph highlights a strong correlation between the returns r_{t+j} and squared returns r_t^2 primarily at positive lags. This pattern indicates that future returns are influenced by current volatility, with most cross-correlation values exceeding the green dashed line, which marks the rejection region for statistical significance. Notably, these correlations primarily occur at positive lags, suggesting that returns are very responsive to past volatility variations.

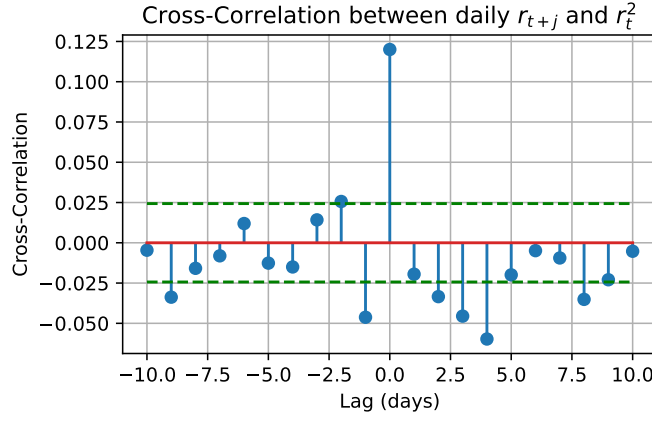


Figure 14: **Empirical cross correlation** of daily lagged log-returns and squared daily returns $\text{corr}(r_{t+j}, r_t^2)$ of AMZN. Sample: **01-21-1999 to 10-16-2024**. The green dotted bars are the (asymptotic) bounds for the rejection region a significance test of each cross-correlation. A line above or below the green dashed line represent a significant cross-correlation.




Figure 15: Time series plot of **AMZN** and **VIX**. Sample: **01-21-1999 to 10-16-2024**.




Figure 16: Scatter plot of daily AMZN log-returns against the daily changes of VIX for the same day t . Fitted OLS regression line is represented in (red). Sample: **01-21-1999 to 10-16-2024**.

The comparison between Amazon's stock and the VIX (Volatility Index) in Figure 15 reveals a negative correlation between the stock and the index. This effect is particularly pronounced during periods of high market volatility, such as the 2008 financial crisis, when the VIX spiked as Amazon's stock experienced a decline. However, this correlation appears less evident during the dot-com bubble, as Amazon maintained strong sales during that period. Overall, the chart reinforces the negative correlation between Amazon's returns and volatility, consistent with the findings from Figure 12.

Figure 16 illustrates the relationship between Amazon's daily log returns and changes in the VIX over the study period. Each blue point represents a specific day, capturing the joint movement of Amazon's log returns and VIX changes. This scatter plot aligns with the findings from Figure 12, highlighting a negative correlation between returns and volatility. The clustering of points around the OLS regression line further supports the presence of a negative correlation between Amazon's returns and its volatility.

A References

References

- [1] Rui Albuquerque. Skewness in stock returns: Reconciling the evidence on firm versus aggregate returns. *The Review of Financial Studies*, 25(5):1630–1673, May 2012. Published: 09 January 2012.
- [2] Giovanni Pumi et al. Estimation of long-range dependent models with missing data: to impute or not to impute? *arXiv preprint*, 2023.

B Python code

Notebook starting next page.

MidTermAssignmentwith8Facts

November 14, 2024

1 Python assignment

```
[86]: # -----  
# File Name: MidTermAssignment.py  
# Description:  
# Autor: Prat Paul; Gavini Charles; Fournier Justin; Blanc Mathieu  
# Creation Date: 2024-10-16  
# Version: 1.0  
# -----
```

Installing yahoofinance

```
[87]: #pip install yfinance
```

Installing statsmodels

```
[88]: #pip install statsmodels
```

```
[89]: #importations  
import numpy as np  
import pandas as pd  
import yfinance as yf  
import matplotlib.pyplot as plt  
import matplotlib.dates as mdates  
from statsmodels.graphics.tsaplots import plot_acf # import this function from  
↳ this submodule  
import statsmodels.api as sm  
import scipy.stats as stats  
from scipy.stats import gaussian_kde, norm, iqr, skew, kurtosis, jarque_bera,↳  
↳ kstest, anderson  
from statsmodels.stats.diagnostic import lilliefors  
import scipy.signal as ss  
import pylab
```

2 First pandas dataframe of Amazon stocks

```
[90]: # Importing Amazon stock from yahoo finance
Amazon = yf.download("AMZN", start="1999-01-21", end="2024-10-16")
Amazon.head()
```

```
[*****100%*****] 1 of 1 completed
```

```
[90]:
```

	Open	High	Low	Close	Adj Close	Volume
Date						
1999-01-21	2.612500	2.759375	2.314063	2.650000	2.650000	940964000
1999-01-22	2.487500	3.146875	2.468750	3.075000	3.075000	875316000
1999-01-25	3.037500	3.084375	2.750000	2.809375	2.809375	546476000
1999-01-26	2.815625	3.031250	2.765625	2.877344	2.877344	490696000
1999-01-27	3.353125	3.493750	3.000000	3.140625	3.140625	700452000

```
[91]: #pip install perfplot
```

```
[92]: latex_table = Amazon.head().to_latex(index=True)
with open("Latex/table.tex", "w") as file:
    file.write(latex_table)
```

```
/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/3304008134.py:1
: FutureWarning: In future versions `DataFrame.to_latex` is expected to utilise
the base implementation of `Styler.to_latex` for formatting and rendering. The
arguments signature may therefore change. It is recommended instead to use
`DataFrame.style.to_latex` which also contains additional functionality.
    latex_table = Amazon.head().to_latex(index=True)
```

3 Cheking if timestamp is 25 years

```
[93]: print('Amazon data range is: ', Amazon.index[0], Amazon.index[-1])

#trying to find gaps

#First create a dataframe for a fullrange of our index, without any gap with
↳ the following formula:
full_range = pd.date_range(start=Amazon.index.min(), end=Amazon.index.max(),
↳ freq='B')

#Then compare to our dataframe:

MissingDays=full_range.difference(Amazon.index)

#Print the count and the detail preview:
print('Missing Days count is: ', len(MissingDays))
print("missing dates", MissingDays)
```

```
print("total size of the 25 years range",len(Amazon.index),"the ratio of_
↳missing inputs/ total size of the data = ",100*len(MissingDays)/len(Amazon.
↳index))
```

#We can see that data have ponctual gaps, no issue here we can still use it

```
Amazon data range is: 1999-01-21 00:00:00 2024-10-15 00:00:00
Missing Days count is: 238
missing dates DatetimeIndex(['1999-02-15', '1999-04-02', '1999-05-31',
'1999-07-05',
                               '1999-09-06', '1999-11-25', '1999-12-24', '2000-01-17',
                               '2000-02-21', '2000-04-21',
                               ...
                               '2023-11-23', '2023-12-25', '2024-01-01', '2024-01-15',
                               '2024-02-19', '2024-03-29', '2024-05-27', '2024-06-19',
                               '2024-07-04', '2024-09-02'],
                             dtype='datetime64[ns]', length=238, freq=None)
total size of the 25 years range 6476 the ratio of missing inputs/ total size of
the data = 3.675108091414453
```

```
[94]: Amazon.index
      #extracting adjusted
      Amzn_adj=Amazon['Adj Close']
      Amzn_adj.index = Amazon.index

      #display first 5 rows, now it is a pandas series instead of a dataframe
      Amzn_adj.head()
```

```
[94]: Date
1999-01-21    2.650000
1999-01-22    3.075000
1999-01-25    2.809375
1999-01-26    2.877344
1999-01-27    3.140625
Name: Adj Close, dtype: float64
```

For the possible gaps in data, we plot them here

```
[95]: #plot the missing dates
full_data = Amazon.reindex(full_range)

#zoom in over one year
start_date = "2010-01-01"
end_date = "2011-01-01"
filtered_full_range = full_range[(full_range >= start_date) & (full_range <=_
↳end_date)]
filtered_missing_dates = MissingDays[(MissingDays >= start_date) & (MissingDays_
↳<= end_date)]
```

```

plt.figure(figsize=(10, 2))

# Plot all dates in the filtered range with gray dots (showing the full
↳ timeline for this period)
plt.plot(filtered_full_range, [1] * len(filtered_full_range), 'o',
↳ color='lightgray', markersize=5, label="Full Date Range")


# Overlay red dots only on the missing dates within the filtered range
plt.plot(filtered_missing_dates, [1] * len(filtered_missing_dates), 'ro',
↳ markersize=5, label="Missing Dates")

# Customize plot
plt.title("Missing Dates in Full Date Range (2010 to 2011)",color='black')
plt.xlabel("Date",color='black')
plt.yticks([]) # Hide y-axis labels for clarity
plt.xticks(rotation=45,color='black')
plt.legend(facecolor='white', edgecolor='black', framealpha=1, fontsize=10)

#Saving the plot in pdf format
plt.savefig('Latex/Img/MissingDates(2010_to_2011).pdf', format='pdf',
↳ bbox_inches='tight')

plt.show()

```



4 PRICES

```

[96]: # extract the closing prices of the Amazon stok (as in lecture)
Pt_d_all = Amazon["Adj Close"]
Pt_d_all.name = 'Pt.d'
# mutate the Index into a DatetimeIndex
Pt_d_all.index = pd.to_datetime(Pt_d_all.index)
Pt_d_all.head()

```

```
[96]: Date
      1999-01-21    2.650000
      1999-01-22    3.075000
      1999-01-25    2.809375
      1999-01-26    2.877344
      1999-01-27    3.140625
      Name: Pt.d, dtype: float64
```

Compute log price

```
[97]: pt_d_all = np.log(Pt_d_all)
      pt_d_all.name = 'pt.d'
      pt_d_all.head()
```

```
[97]: Date
      1999-01-21    0.974560
      1999-01-22    1.123305
      1999-01-25    1.032962
      1999-01-26    1.056868
      1999-01-27    1.144422
      Name: pt.d, dtype: float64
```

Compute weekly monthly and yearly

```
[98]: pt_w_all = pt_d_all.resample('W').last()
      pt_m_all = pt_d_all.resample('M').last()
      pt_y_all = pt_d_all.resample('Y').last()
      # and rename them:
      pt_w_all.name = 'pt.w.all'
      pt_m_all.name = 'pt.m.all'
      pt_y_all.name = 'pt.y.all'

      #idem for simply prices
      Pt_w_all = Pt_d_all.resample('W').last()
      Pt_m_all = Pt_d_all.resample('M').last()
      Pt_y_all = Pt_d_all.resample('Y').last()
      # and rename them:
      Pt_w_all.name = 'Pt_w_all'
      Pt_m_all.name = 'Pt_m_all'
      Pt_y_all.name = 'Pt_y_all'
```

Plot the simple prices

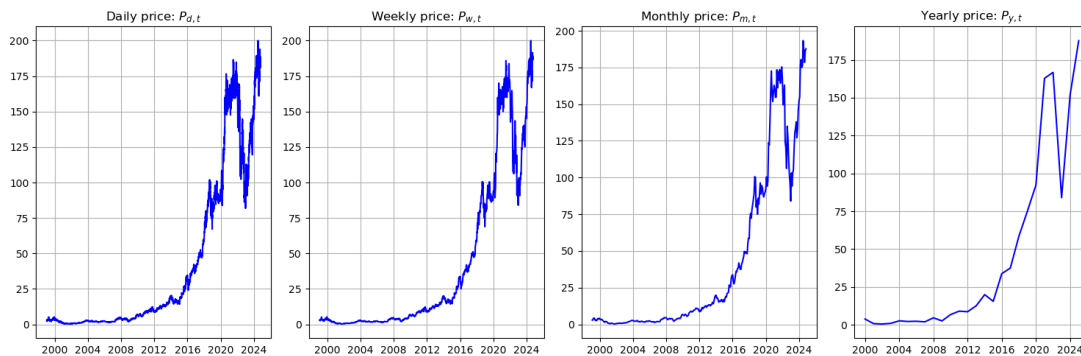
```
[99]: # set the 1x4 windows layout
      fig, axs = plt.subplots(1, 4, figsize=(15, 5))
      # Daily Price
      axs[0].plot(Pt_d_all.index, Pt_d_all, color='blue')
      axs[0].set_title('Daily price: $P_{d,t}$')
      axs[0].grid(True)
```

```

# Weekly price
axs[1].plot(Pt_w_all.index, Pt_w_all, color='blue')
axs[1].set_title('Weekly price:  $P_{w,t}$ ')
axs[1].grid(True)
# Monthly price
axs[2].plot(Pt_m_all.index, Pt_m_all, color='blue')
axs[2].set_title('Monthly price:  $P_{m,t}$ ')
axs[2].grid(True)
# Yearly price
axs[3].plot(Pt_y_all.index, Pt_y_all, color='blue')
axs[3].set_title('Yearly price:  $P_{y,t}$ ')
axs[3].grid(True)

# Manage margins and plot
plt.tight_layout()
plt.savefig('Latex/Img/prices_time.pdf', format='pdf', bbox_inches='tight')
plt.show()

```



Adding python code to the latex document in the appendix part

```

[100]: #Test for incorporating python code into the appendix section in the latex_
        ↪ document
code_content = r"""
\section{Appendix: Python Code}
Below is the Python code used in this analysis.

\begin{lstlisting}[language=Python, caption=Python Code for Analysis]
# Python code example
import numpy as np
import pandas as pd

def analyze_data(data):
    mean = np.mean(data)
    std_dev = np.std(data)

```



```

    return mean, std_dev

data = [1, 2, 3, 4, 5]
mean, std_dev = analyze_data(data)
print(f"Mean: {mean}, Standard Deviation: {std_dev}")
\end{lstlisting}
"""

# Write to the 'code_appendix.tex' file
with open("Latex/code_appendix.tex", "w") as file:
    file.write(code_content)

```

5 *Calculating returns*

```

[101]: #calculating return

#log returns VS simple returns
Rt_d_all_temp = Pt_d_all.pct_change()
rt_d_all_temp = pt_d_all.diff()
rt_d_all_temp, Rt_d_all_temp

```

```

[101]: (Date
1999-01-21      NaN
1999-01-22    0.148745
1999-01-25   -0.090343
1999-01-26    0.023906
1999-01-27    0.087554
...
2024-10-09    0.013319
2024-10-10    0.007961
2024-10-11    0.011559
2024-10-14   -0.006802
2024-10-15    0.000800
Name: pt.d, Length: 6476, dtype: float64,
Date
1999-01-21      NaN
1999-01-22    0.160377
1999-01-25   -0.086382
1999-01-26    0.024194
1999-01-27    0.091501
...
2024-10-09    0.013408
2024-10-10    0.007993
2024-10-11    0.011626
2024-10-14   -0.006779
2024-10-15    0.000800

```

Name: Pt.d, Length: 6476, dtype: float64)

Compute daily, weekly, and monthly

```
[102]: rt_d_all = pt_d_all.diff().dropna() #dropna remove the first NaN
rt_w_all = pt_w_all.diff().dropna()
rt_m_all = pt_m_all.diff().dropna()
rt_y_all = pt_y_all.diff().dropna()

Rt_d_all = Pt_d_all.pct_change().dropna() #dropna remove the first NaN
Rt_w_all = Pt_w_all.pct_change().dropna()
Rt_m_all = Pt_m_all.pct_change().dropna()
Rt_y_all = Pt_y_all.pct_change().dropna()

# and rename them:
rt_d_all.name = 'rt_d_all'
rt_w_all.name = 'rt_w_all'
rt_m_all.name = 'rt_m_all'
rt_y_all.name = 'rt_y_all'

Rt_d_all.name = 'Rt_d_all'
Rt_w_all.name = 'Rt_w_all'
Rt_m_all.name = 'Rt_m_all'
Rt_y_all.name = 'Rt_y_all'

rt_d_all.head()
Rt_d_all.head()
```

```
[102]: Date
1999-01-22    0.160377
1999-01-25   -0.086382
1999-01-26    0.024194
1999-01-27    0.091501
1999-01-28   -0.021891
Name: Rt_d_all, dtype: float64
```

The first returns are correctly computed, we have to be careful to the dropna

Let's plot returns

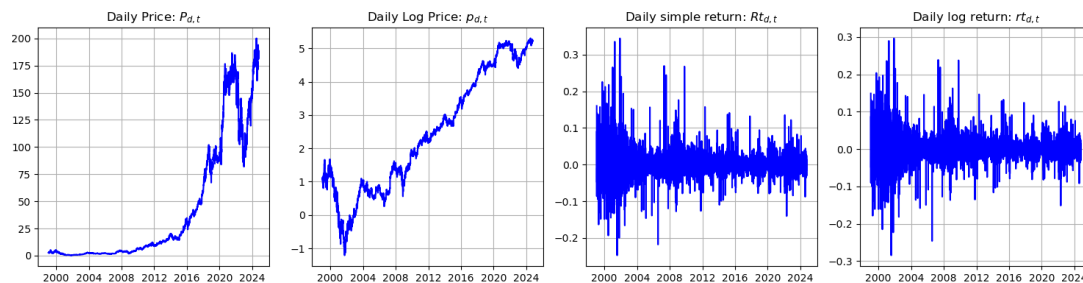
```
[103]: # set the 1x3 windows layout
fig, axs = plt.subplots(1, 4, figsize=(15, 4))
# Daily Price
axs[0].plot(Pt_d_all.index, Pt_d_all, color='blue')
axs[0].set_title('Daily Price: $P_{d,t}$')
axs[0].grid(True)
# Daily log price
axs[1].plot(pt_d_all.index, pt_d_all, color='blue')
axs[1].set_title('Daily Log Price: $p_{d,t}$')
```

```

axs[1].grid(True)
# Daily simple returns
axs[2].plot(Rt_d_all.index, Rt_d_all, color='blue')
axs[2].set_title('Daily simple return:  $R_{t,d,t}$ ')
axs[2].grid(True)
# Daily log returns
axs[3].plot(rt_d_all.index, rt_d_all, color='blue')
axs[3].set_title('Daily log return:  $rt_{d,t}$ ')
axs[3].grid(True)

plt.tight_layout()
plt.savefig('Latex/Img/log_returns.pdf', format='pdf', bbox_inches='tight')
plt.show()

```



Squared returns

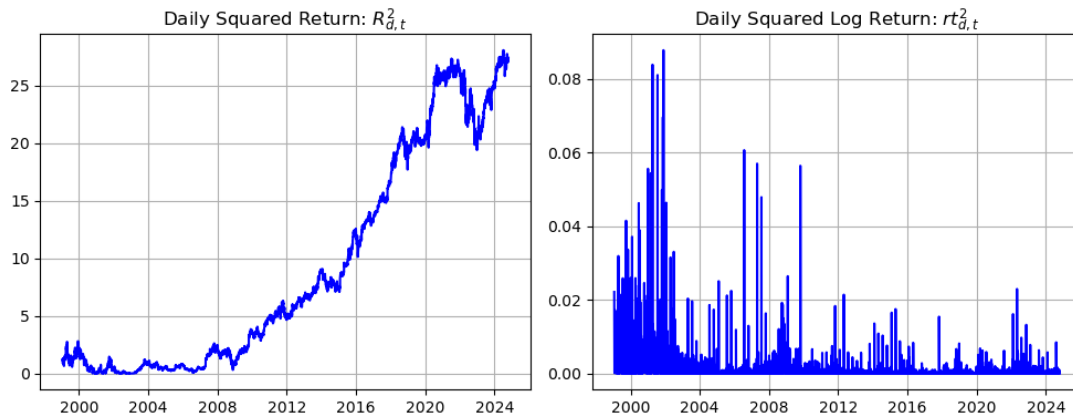
```

[104]: fig, axs = plt.subplots(1, 2, figsize=(10, 4))

# Daily squared log price
axs[0].plot(pt_d_all.index, pt_d_all**2, color='blue')
axs[0].set_title('Daily Squared Return:  $R_{t,d,t}^2$ ')
axs[0].grid(True)
# Daily squared log returns
axs[1].plot(rt_d_all.index, rt_d_all**2, color='blue')
axs[1].set_title('Daily Squared Log Return:  $rt_{d,t}^2$ ')
axs[1].grid(True)

plt.tight_layout()
plt.savefig('Latex/Img/squared_log_returns.pdf', format='pdf',
    ↪bbox_inches='tight')
plt.show()

```

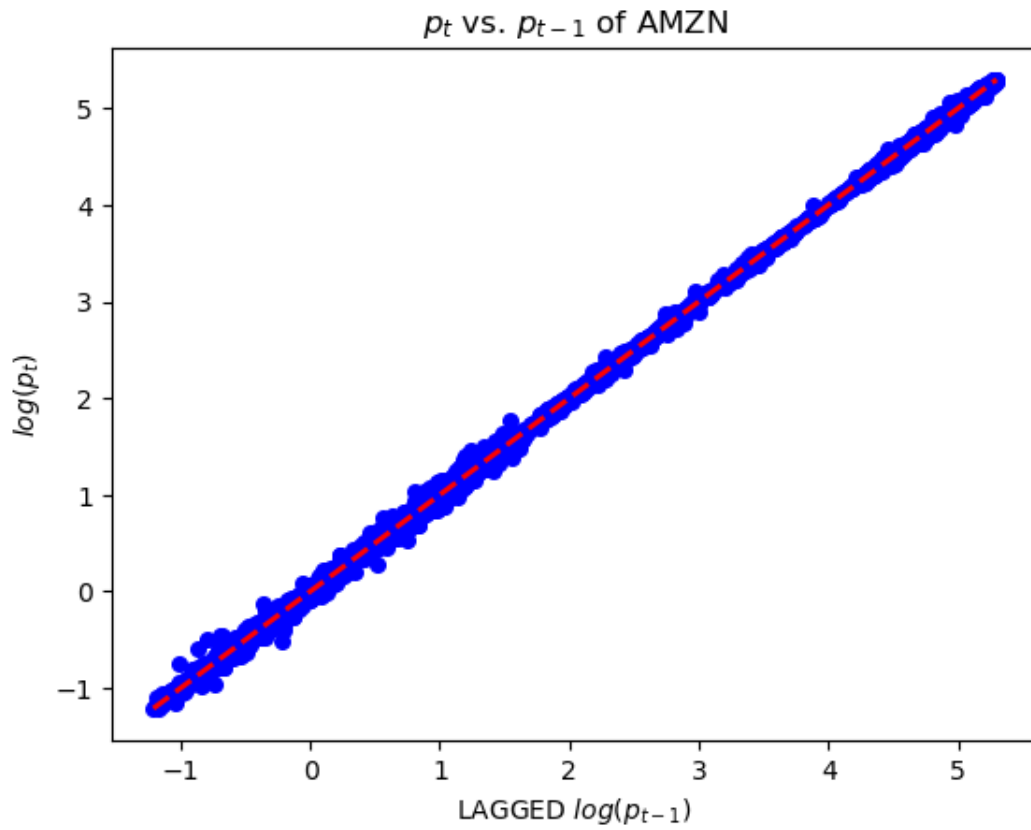


6 Scatterplot of $p_t - p_{t-1}$

First define the function for plotting a scatterplot

```
[105]: def lag1_scatterplot(data, x_label, y_label, title):
    plt.scatter(data.shift(), data, color='blue', s=30) #data.shift is just
    ↪ p_{t-1} (if t=0 what happend ?)
    plt.xlabel(x_label)
    plt.ylabel(y_label)
    plt.title(title)
    plt.plot([min(data), max(data)], [min(data), max(data)],
    ↪ linestyle='dashed', linewidth=2, color='red')
    # plt.savefig('Latex/Img/Laggedlog(p_{t-1}).pdf', format='pdf',
    ↪ bbox_inches='tight')

[106]: lag1_scatterplot(pt_d_all, "LAGGED $log(p_{t-1})$", "$log(p_t)$", "$p_t$ vs.
    ↪ $p_{t-1}$ of AMZN")
```



7 Autocorrelation

```
[107]: """autocorrelate=pt_d_all.shift().corrwith(pt_d_all, method='pearson')
print(round(autocorrelate,4))"""
```

```
[107]: "autocorrelate=pt_d_all.shift().corrwith(pt_d_all,
method='pearson')\nprint(round(autocorrelate,4))"
```

```
[108]: """
autocorrelate = pt_d_all.shift(1).corrwith(pt_d_all, method='pearson')
print(autocorrelate.round(4))
"""
```

```
[108]: "\nautocorrelate = pt_d_all.shift(1).corrwith(pt_d_all,
method='pearson')\nprint(autocorrelate.round(4))\n"
```

8 4.1/ Prices are non-stationary

1. Profile of Log Prices with Time

2. P_t VS P_{t-1}

3. Autocorrelation of Daily Prices

```
[109]: import matplotlib.pyplot as plt

# Set the layout for 1x3 subplots (though you may not need all subplots)
fig, axs = plt.subplots(1, 1, figsize=(15, 4))

# Plot Daily Price
axs.plot(pt_d_all.index, pt_d_all, color='blue')
axs.set_title('Log Daily Price:  $p_{d,t}$ ')
axs.grid(True)

# Show the plot
plt.show()
```



```
[110]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np

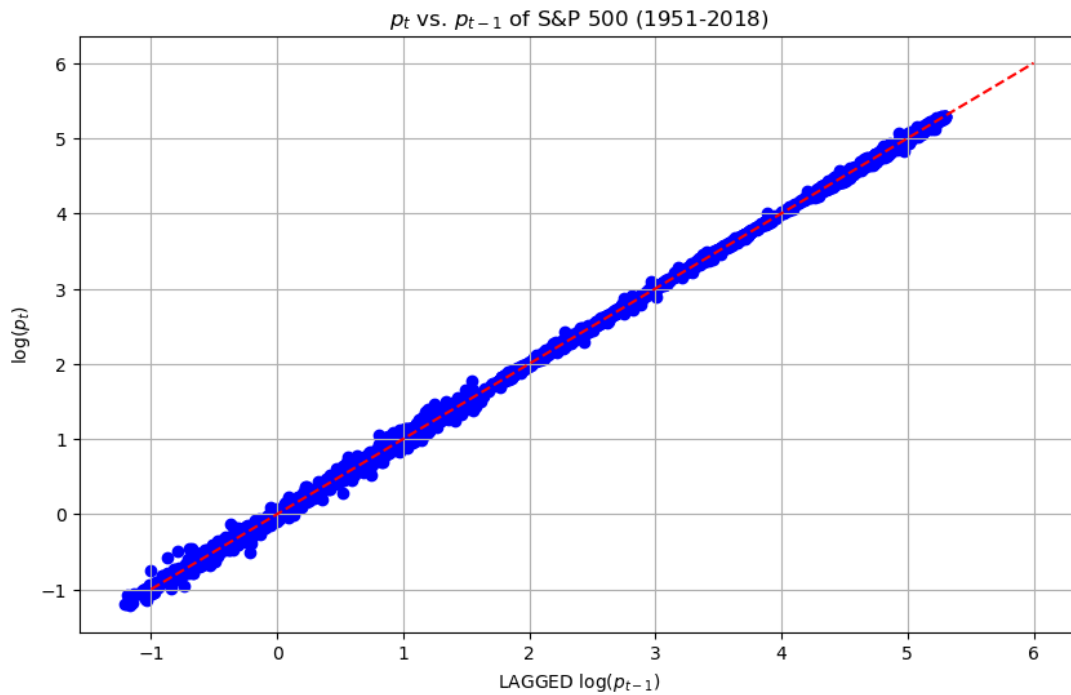
# Estrai i log-prezzi giornalieri
log_price_daily = pt_d_all

# Calcola il log-prezzo al giorno precedente
log_price_previous = log_price_daily.shift(1)

# Creazione dello scatter plot
plt.figure(figsize=(10, 6))
plt.scatter(log_price_previous, log_price_daily, color='blue')
plt.plot([-1, 6], [-1, 6], color='red', linestyle='--')
plt.title('$p_t$ vs.  $p_{t-1}$  of S&P 500 (1951-2018)')
plt.xlabel(r'LAGGED  $\log(p_{t-1})$ ')
plt.ylabel(r' $\log(p_t)$ ')
plt.grid(True)
```

```
# Salva il grafico in formato png
plt.savefig('Latex/Img/log(pt) vs log(pt-1).pdf', format='pdf',
            bbox_inches='tight')

plt.show()
```



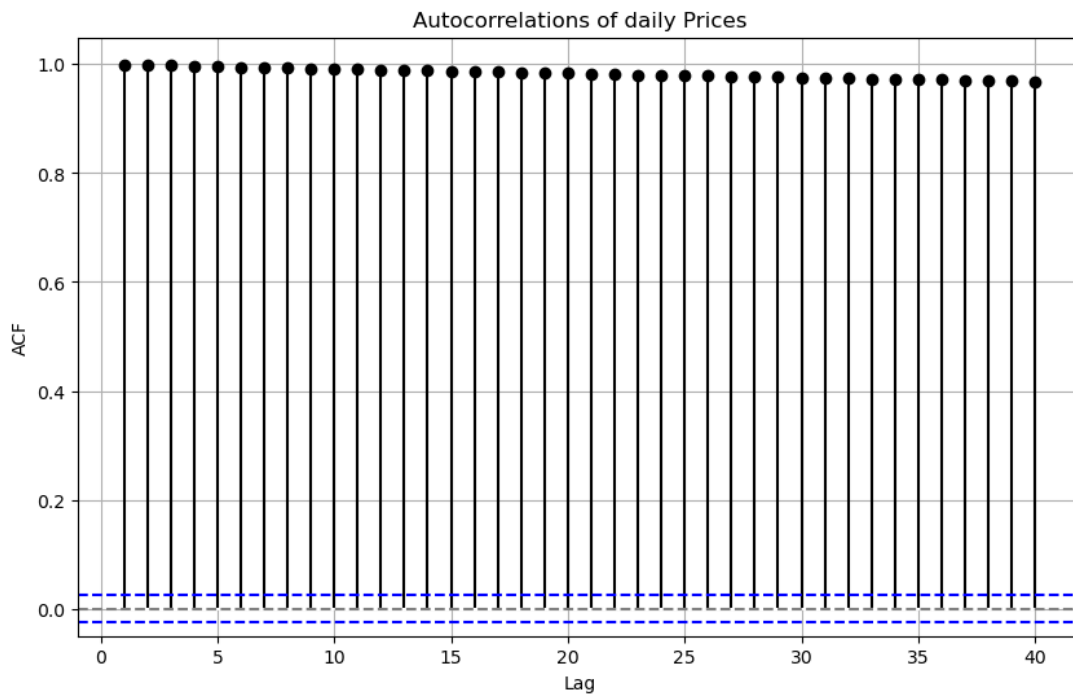
```
[111]: from statsmodels.tsa.stattools import acf

# Calculate empirical autocorrelation
lags = 40
acf_values = acf(Pt_d_all, nlags=lags)

# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(Pt_d_all))

# Create the autocorrelation plot with Bartlett intervals
plt.figure(figsize=(10, 6))
plt.stem(np.arange(1, lags + 1), acf_values[1:], linefmt='k-', markerfmt='ko',
         basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
```

```
plt.title('Autocorrelations of daily Prices')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)
#plt.savefig('Latex/Autocorrel_daily.pdf', format='pdf', bbox_inches='tight')
plt.show()
```



9 4.2/ Log Returns are Stationary

1. Profile of Log Returns with Time
2. R_t VS $R(t-1)$
3. Autocorrelation of Daily Returns

```
[112]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np

# Calculate daily log returns
log_returns_daily = rt_d_all

# Create the plot of daily log returns with a black horizontal line
plt.figure(figsize=(10, 6))
```



```

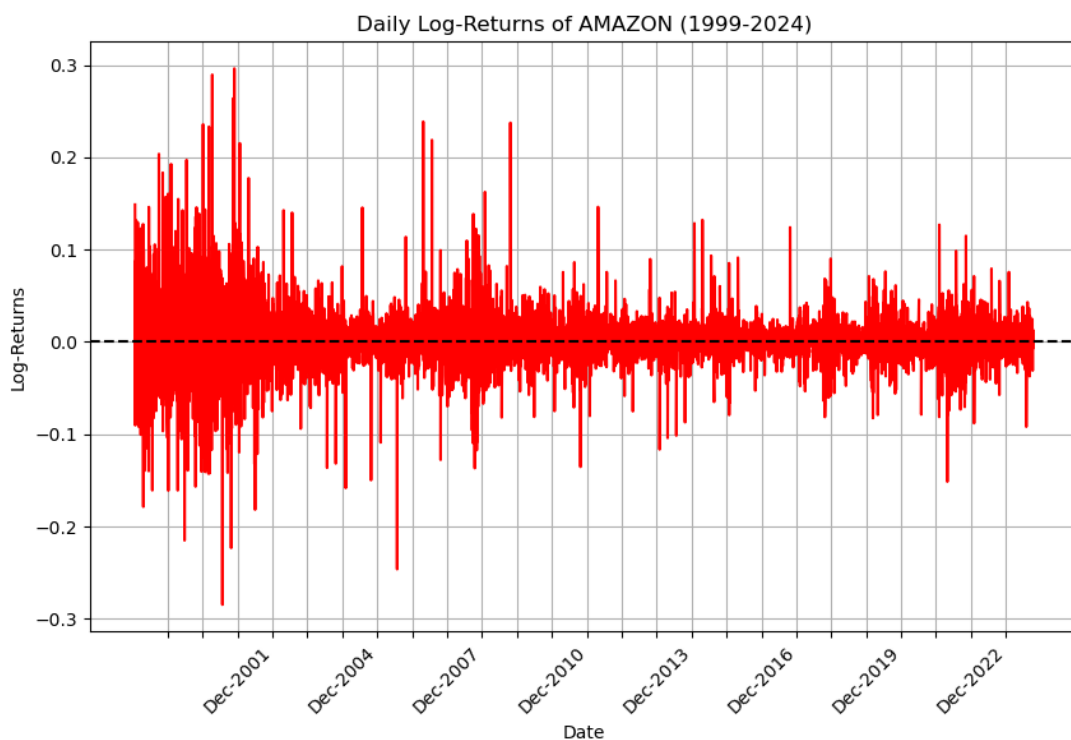
plt.plot(log_returns_daily.index, log_returns_daily, color='red')
plt.axhline(y=0, color='black', linestyle='--')
plt.title('Daily Log>Returns of AMAZON (1999-2024)')
plt.xlabel('Date')
plt.ylabel('Log>Returns')
plt.grid(True)

# Customizing x-axis labels for December 31 of each year
date_labels = pd.date_range(start='1999-12-31', end='2023-12-31', freq='A-DEC')
# Show 1 tick every 3 years
formatted_labels = [f'Dec-{date.year}' if date.year % 3 == 0 else '' for date_
    ↪in date_labels]
# Add labels and rotate them
plt.xticks(date_labels, formatted_labels, rotation=45)

# Save the plot in png format
plt.savefig('Latex/Img/Daily Log Returns.pdf', format='pdf',
    ↪bbox_inches='tight')

plt.show()

```



```
[113]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np

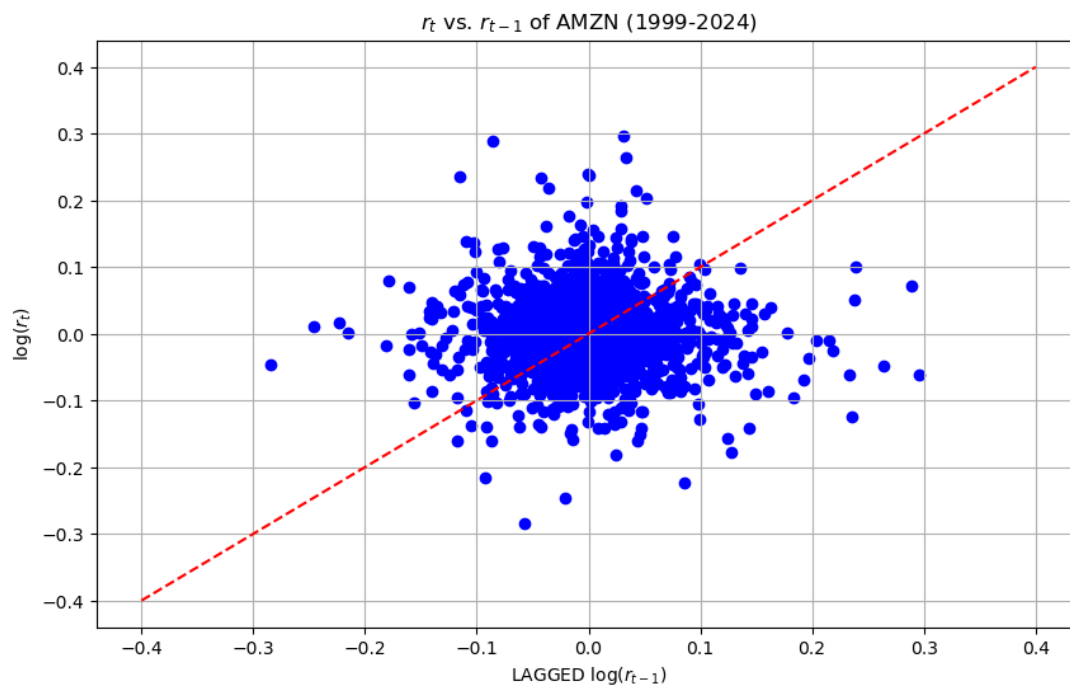
# Get the Daily Log Returns
log_return_daily = rt_d_all

# Calculation of the Lagged log returns
log_return_previous = log_return_daily.shift(1)

# Creation of the Scatter Plot
plt.figure(figsize=(10, 6))
plt.scatter(log_return_previous, log_return_daily, color='blue')
plt.plot([-0.4, 0.4], [-0.4, 0.4], color='red', linestyle='--')
plt.title('$r_t$ vs. $r_{t-1}$ of AMZN (1999-2024)')
plt.xlabel(r'$\text{LAGGED } \log(r_{t-1})$')
plt.ylabel(r'$\log(r_t)$')
plt.grid(True)

# Saving the Image
plt.savefig('Latex/Img/LogReturns_vs_LaggedLogReturns.pdf', format='png',
           bbox_inches='tight')

plt.show()
```

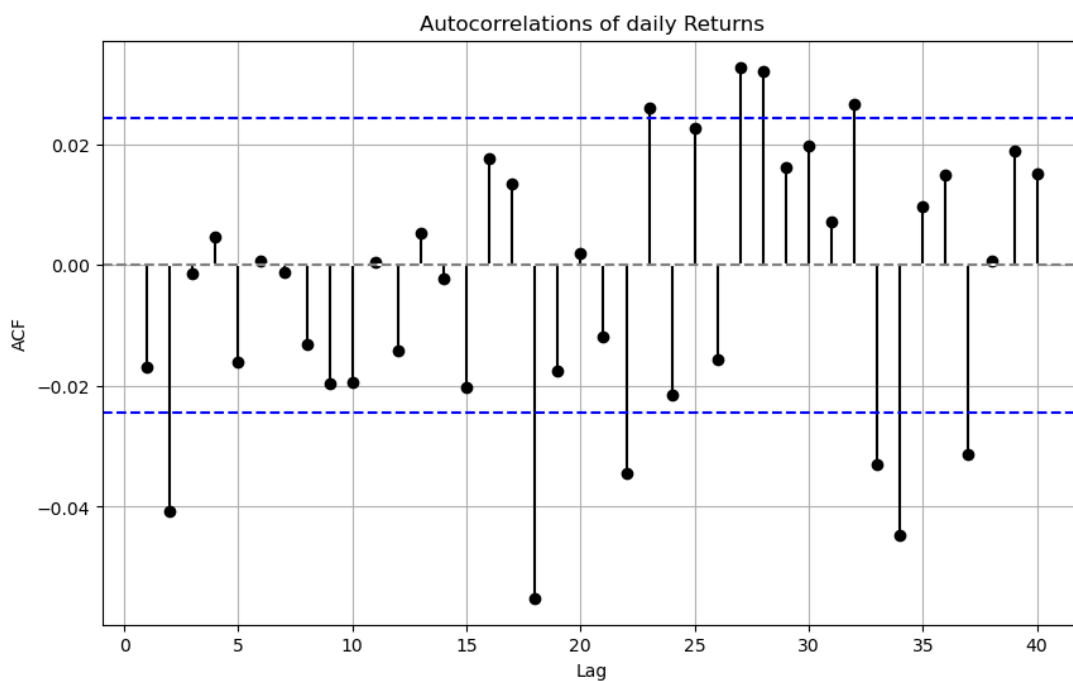


```
[114]: from statsmodels.tsa.stattools import acf

# Calculate empirical autocorrelation
lags = 40
acf_values = acf(Rt_d_all, nlags=lags)

# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(Rt_d_all))

# Create the autocorrelation plot with Bartlett intervals
plt.figure(figsize=(10, 6))
plt.stem(np.arange(1, lags + 1), acf_values[1:], linefmt='k-', markerfmt='ko',
        basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
plt.title('Autocorrelations of daily Returns')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)
#plt.savefig('Latex/Autocorrel_Returns_daily.pdf', format='pdf',
#        bbox_inches='tight')
plt.show()
```



10 4.3/ Are Log Returns Asymmetric ?

1. Rolling Mean
2. Rolling Standard Deviation
3. Rolling Skewness
4. Current Skewness and Interpretation

```
[161]: import yfinance as yf
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import skew, kurtosis
import pandas as pd

# Compute daily log-returns
log_returns_daily = rt_d_all

# set the rolling window equal to 252 days
window_length = 252
T = log_returns_daily.shape[0]

# Create an empty matrix to store data
roll_mom_manual = np.zeros((T, 5))

# Run a for loop to fill the matrix with moments
for i in range(window_length, T):
    est_window = np.arange(i - window_length + 1, i + 1)

    # Use .iloc to select rows by integer positions, not labels
    y = log_returns_daily.iloc[est_window]

    # Compute the moments for each
    roll_mom_manual[i, 0] = np.mean(y)
    roll_mom_manual[i, 1] = np.std(y, ddof=1)
    roll_mom_manual[i, 2] = skew(y)
    roll_mom_manual[i, 3] = kurtosis(y)
    roll_mom_manual[i, 4] = np.mean((y - np.mean(y))**4)

# Plot results of manually computed rolling mean
mean_plot_man = roll_mom_manual[:, 0]
mean_plot_man_ub = mean_plot_man + 1.96 * roll_mom_manual[:, 1] / np.
↳sqrt(window_length)
mean_plot_man_lb = mean_plot_man - 1.96 * roll_mom_manual[:, 1] / np.
↳sqrt(window_length)
```

```

data2plot_na = np.column_stack((mean_plot_man, mean_plot_man_lb,
    ↪mean_plot_man_ub))

data_index = log_returns_daily.index

data2plot_na = pd.DataFrame({'Mean': mean_plot_man, 'LowerBound':
    ↪mean_plot_man_lb, 'UpperBound': mean_plot_man_ub},
                             index=data_index)

# Select only rows without missing values
data2plot = data2plot_na.dropna()
# retrieve the data index
data2plot

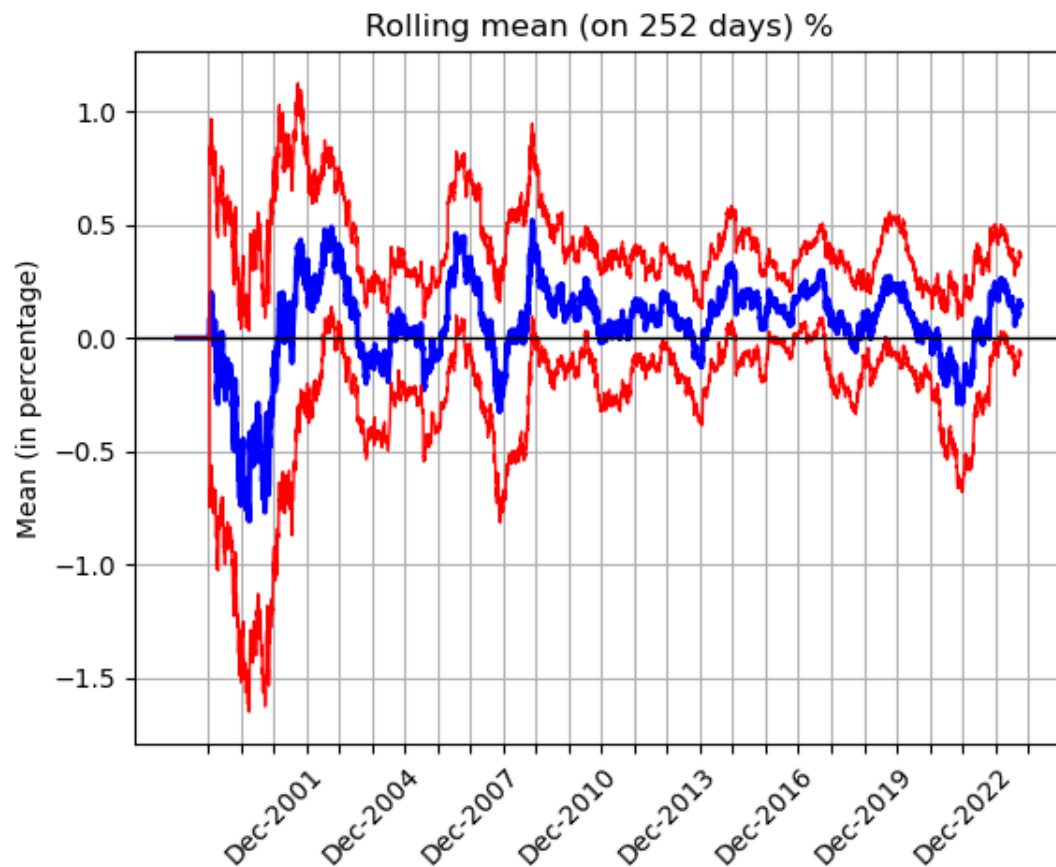
# Customizing x-axis labels for December 31 of each year
date_labels = pd.date_range(start='1999-12-31', end='2024-12-31', freq='A-DEC')
# Show 1 tick every 3 years
formatted_labels = [f'Dec-{date.year}' if date.year % 3 == 0 else '' for date
    ↪in date_labels]
# Add labels and rotate them
plt.xticks(date_labels, formatted_labels, rotation=45)

# Plot the data
plt.plot(data2plot.index, data2plot["Mean"] * 100, color='blue', linestyle='-',
    ↪linewidth=2)
plt.plot(data2plot.index, data2plot["LowerBound"] * 100, color='red',
    ↪linestyle='-', linewidth=1)
plt.plot(data2plot.index, data2plot["UpperBound"] * 100, color='red',
    ↪linestyle='-', linewidth=1)
plt.grid(True)
plt.xlabel('')
plt.ylabel('Mean (in percentage)')
plt.title('Rolling mean (on 252 days) %')
plt.axhline(0, linestyle='-', color='black', linewidth=1) # Add a zero line

plt.savefig('Latex/Img/AMZN_MEAN_rolling_1999_2024.pdf', format='pdf',
    ↪bbox_inches='tight')

plt.show()

```



```
[162]: # extract the Std Dev from roll_mom_manual
sd_plot = roll_mom_manual[:,1]
mu4 = roll_mom_manual[:,4]
sd_plot_ub = roll_mom_manual[:,1]+1.96*(1/(2*sd_plot)*np.sqrt(mu4-sd_plot**4))/
↳ np.sqrt(window_length)
sd_plot_lb = roll_mom_manual[:,1]-1.96*(1/(2*sd_plot)*np.sqrt(mu4-sd_plot**4))/
↳ np.sqrt(window_length)

data2plot_na = np.column_stack((sd_plot, sd_plot_lb, sd_plot_ub))

data_index = log_returns_daily.index

data2plot_na = pd.DataFrame({'Std': sd_plot, 'LowerBound': sd_plot_lb,
↳ 'UpperBound': sd_plot_ub},
                             index=data_index)

# Select only rows without missing values
data2plot = data2plot_na.dropna()
```

```

# retrieve the data index
data2plot

# Customizing x-axis labels for December 31 of each year
date_labels = pd.date_range(start='1999-01-01', end='2024-10-01', freq='A-DEC')
# Show 1 tick every 3 years
formatted_labels = [f'Dec-{date.year}' if date.year % 3 == 0 else '' for date
    ↪in date_labels]
# Add labels and rotate them
plt.xticks(date_labels, formatted_labels, rotation=45)

# Plot the data
plt.plot(data2plot.index, data2plot["StD"] * 100, color='blue', linestyle='-',
    ↪linewidth=2)
plt.plot(data2plot.index, data2plot["LowerBound"] * 100, color='red',
    ↪linestyle='-', linewidth=1)
plt.plot(data2plot.index, data2plot["UpperBound"] * 100, color='red',
    ↪linestyle='-', linewidth=1)
plt.xlabel('')
plt.grid(True)
plt.ylabel('st.dev. (%)')
plt.title('Rolling st.dev. (on 252 days) %')
plt.axhline(0, linestyle='-', color='black', linewidth=1) # Add a zero line

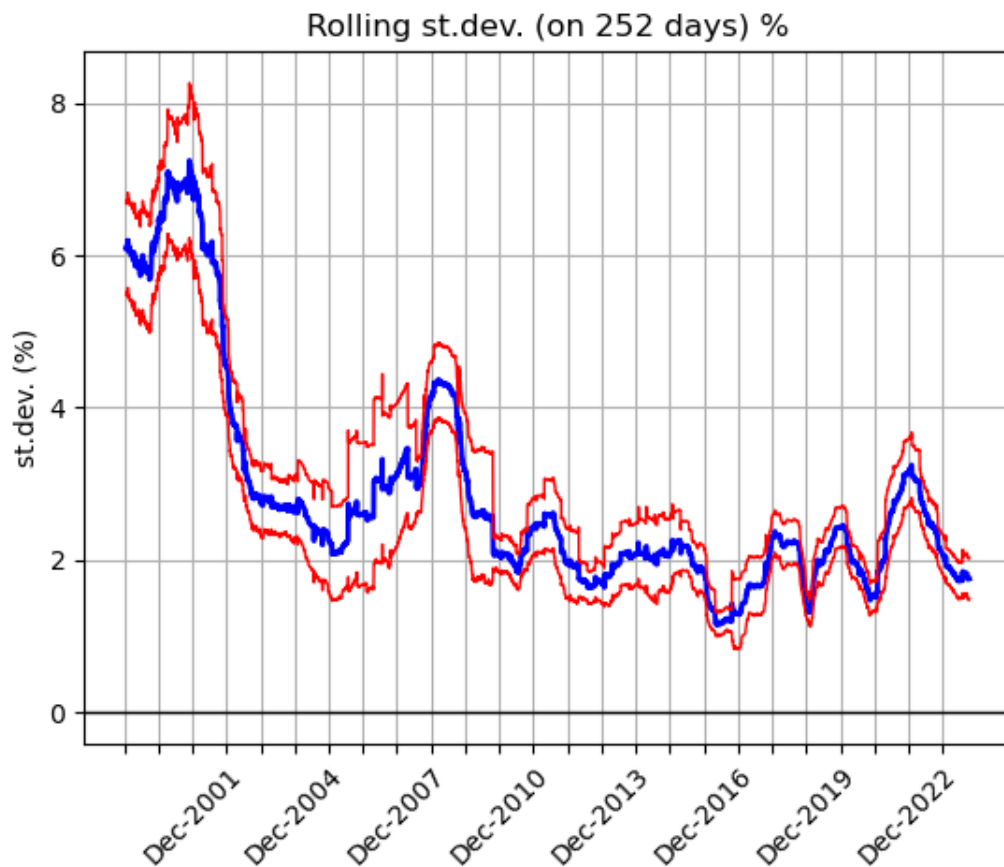
plt.savefig('Latex/Img/Fact7_AMZN_rolling_stdev.pdf', format='pdf',
    ↪bbox_inches='tight')
plt.show()

```

```

/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/1880708915.py:4
: RuntimeWarning: divide by zero encountered in true_divide
    sd_plot_ub = roll_mom_manual[:,1]+1.96*(1/(2*sd_plot)*np.sqrt(mu4-
sd_plot**4))/np.sqrt(window_length)
/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/1880708915.py:4
: RuntimeWarning: invalid value encountered in multiply
    sd_plot_ub = roll_mom_manual[:,1]+1.96*(1/(2*sd_plot)*np.sqrt(mu4-
sd_plot**4))/np.sqrt(window_length)
/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/1880708915.py:5
: RuntimeWarning: divide by zero encountered in true_divide
    sd_plot_lb = roll_mom_manual[:,1]-1.96*(1/(2*sd_plot)*np.sqrt(mu4-
sd_plot**4))/np.sqrt(window_length)
/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/1880708915.py:5
: RuntimeWarning: invalid value encountered in multiply
    sd_plot_lb = roll_mom_manual[:,1]-1.96*(1/(2*sd_plot)*np.sqrt(mu4-
sd_plot**4))/np.sqrt(window_length)

```



```
[116]: # Skewness
skew_plot = roll_mom_manual[:,2]
skew_plot_ub = np.full(skew_plot.shape[0],+1.96*np.sqrt(6)/np.
    ↪sqrt(window_length))
skew_plot_lb = np.full(skew_plot.shape[0],-1.96*np.sqrt(6)/np.
    ↪sqrt(window_length))

data2plot_na = np.column_stack((skew_plot, skew_plot_lb, skew_plot_ub))

data_index = log_returns_daily.index

data2plot_na = pd.DataFrame({'Skewness': skew_plot, 'LowerBound': skew_plot_lb,
    ↪'UpperBound': skew_plot_ub},
    index=data_index)

# Select only rows without missing values
data2plot = data2plot_na.dropna()
# retrieve the data index
```



```

data2plot

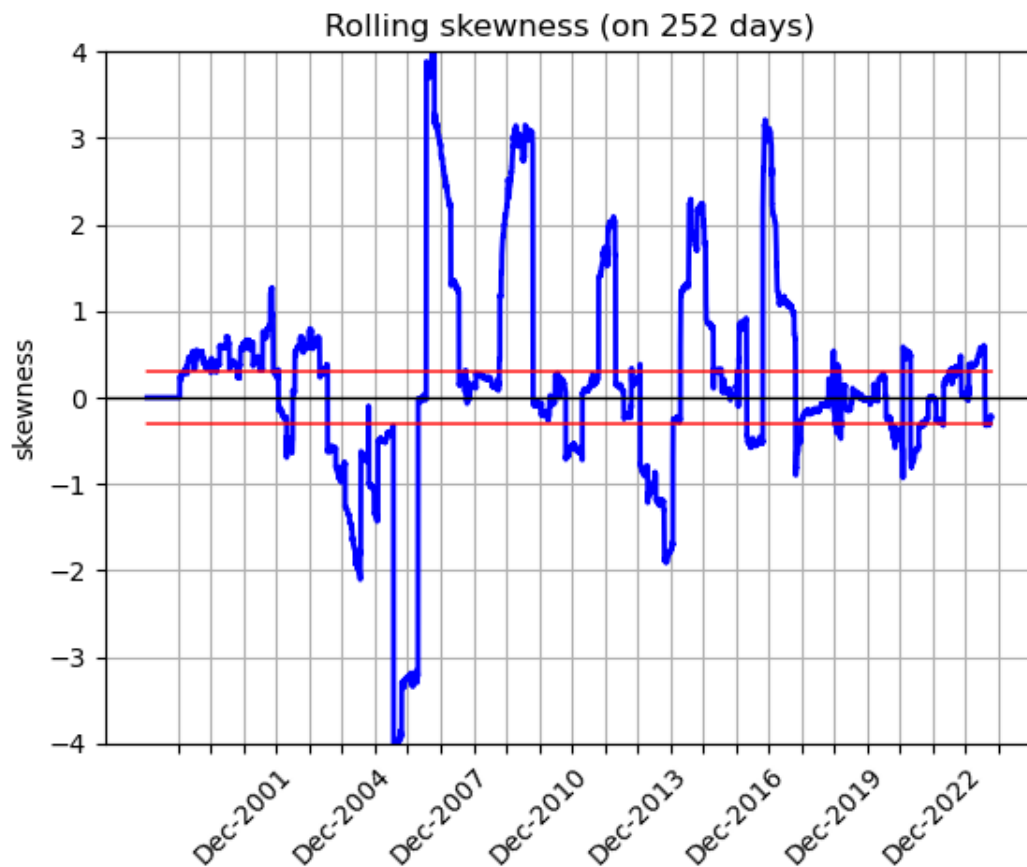
# Customizing x-axis labels for December 31 of each year
date_labels = pd.date_range(start='1999-12-31', end='2024-12-31', freq='A-DEC')
# Show 1 tick every 3 years
formatted_labels = [f'Dec-{date.year}' if date.year % 3 == 0 else '' for date_
    ↪in date_labels]
# Add labels and rotate them
plt.xticks(date_labels, formatted_labels, rotation=45)

# Plot the data
plt.plot(data2plot.index, data2plot["Skewness"], color='blue', linestyle='-',
    ↪linewidth=2)
plt.plot(data2plot.index, data2plot["LowerBound"], color='red', linestyle='-',
    ↪linewidth=1)
plt.plot(data2plot.index, data2plot["UpperBound"], color='red', linestyle='-',
    ↪linewidth=1)
plt.ylim(-4,4)
plt.grid(True)
plt.xlabel('')
plt.ylabel('skewness')
plt.title('Rolling skewness (on 252 days)')
plt.axhline(0, linestyle='-', color='black', linewidth=1) # Add a zero line

plt.savefig('Latex/Img/AMZN_skew_rolling_1999_2024.pdf', format='pdf',
    ↪bbox_inches='tight')

plt.show()

```



11 4.4/ *Heavy Tailed Distribution for the Daily Log Returns ?*

1. Comparison of the Normal Distribution vs our actual Values
2. Excess Kurtosis of our data
3. Interpretation

```
[117]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats as stats
import seaborn as sns

# Extract daily log-returns
log_price_daily = pt_d_all # Ensure this is a pandas DataFrame or Series
log_returns_daily = rt_d_all # Ensure this is a pandas DataFrame or Series
```

```

# If log_returns_daily is a DataFrame, convert it to a 1D array (assuming
↳ 'column_name' is the name of the column)
log_returns_daily = log_returns_daily.values.flatten() # Ensure it's 1D array

# Calculate monthly log-returns
log_price_weekly = pt_w_all # Ensure this is a pandas DataFrame or Series
log_returns_weekly = rt_w_all # Ensure this is a pandas DataFrame or Series

# If log_returns_monthly is a DataFrame, convert it to a 1D array (assuming
↳ 'column_name' is the name of the column)
log_returns_weekly = log_returns_weekly.values.flatten() # Ensure it's 1D array

# Create the figure with four subplots
fig, axs = plt.subplots(2, 2, figsize=(12, 8))

# Plot histogram of daily log-returns
sns.histplot(log_returns_daily, bins=30, color='lime', edgecolor='black',
↳ kde_kws={'color': 'red'}, ax=axs[0, 0], stat='density')
axs[0, 0].plot(np.linspace(log_returns_daily.min(), log_returns_daily.max(),
↳ 100),
               stats.norm.pdf(np.linspace(log_returns_daily.min(),
↳ log_returns_daily.max(), 100),
                             log_returns_daily.mean(), log_returns_daily.
↳ std()), color='red', linewidth=2)
axs[0, 0].set_title('Histogram of Daily Log>Returns')
axs[0, 0].set_xlabel('Log>Returns')
axs[0, 0].set_ylabel('Density')

# Plot histogram of monthly log-returns
sns.histplot(log_returns_weekly, bins=30, color='lime', edgecolor='black',
↳ kde_kws={'color': 'red'}, ax=axs[0, 1], stat='density')
axs[0, 1].plot(np.linspace(log_returns_weekly.min(), log_returns_weekly.max(),
↳ 100),
               stats.norm.pdf(np.linspace(log_returns_weekly.min(),
↳ log_returns_weekly.max(), 100),
                             log_returns_weekly.mean(), log_returns_weekly.
↳ std()), color='red', linewidth=2)
axs[0, 1].set_title('Histogram of Weekly Log>Returns')
axs[0, 1].set_xlabel('Log>Returns')
axs[0, 1].set_ylabel('Density')

# QQ plot of daily log-returns
stats.probplot(log_returns_daily, dist="norm", plot=axs[1, 0])
axs[1, 0].set_title('QQ Plot of Daily Log>Returns')
axs[1, 0].set_xlabel('Normal Quantiles')
axs[1, 0].set_ylabel('Sample Quantiles')

```

```

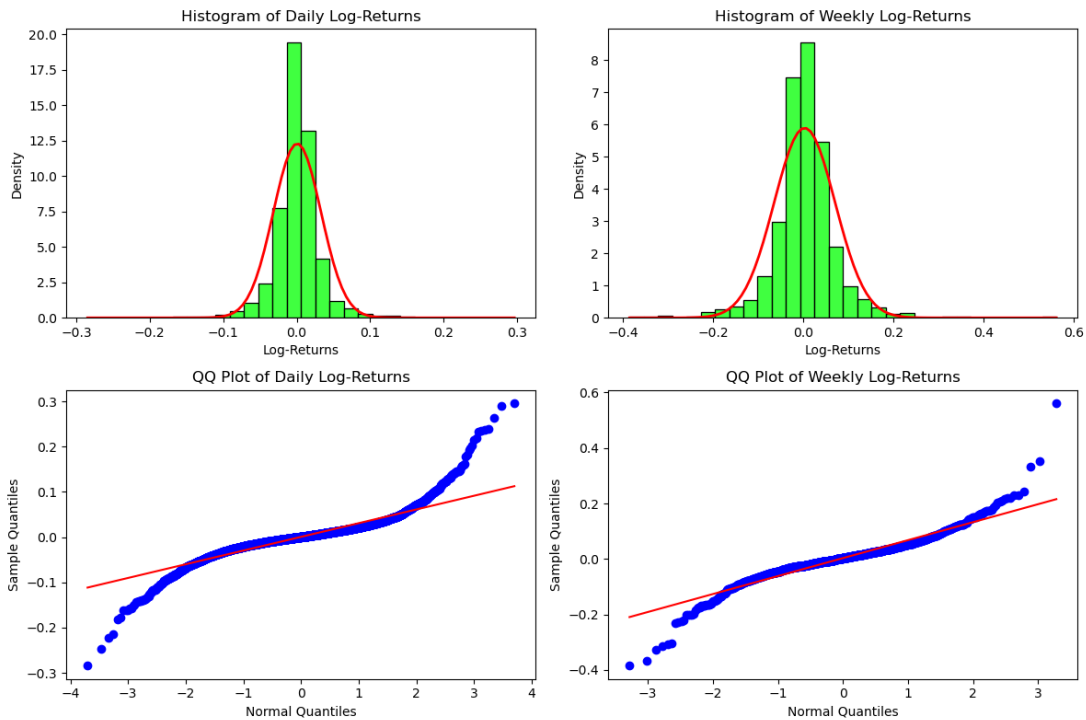
# QQ plot of monthly log-returns
stats.probplot(log_returns_weekly, dist="norm", plot=axes[1, 1])
axes[1, 1].set_title('QQ Plot of Weekly Log>Returns')
axes[1, 1].set_xlabel('Normal Quantiles')
axes[1, 1].set_ylabel('Sample Quantiles')

# Adjust spacing between plots
plt.tight_layout()

# Save the plot in pdf format
plt.savefig('Latex/Img/QQplot_daily_weekly_AMZN.pdf', format='pdf',
            bbox_inches='tight')

# Show the plot
plt.show()

```



```

[118]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats as stats

# Extract daily log-returns

```

```

log_returns_daily = rt_d_all.values.flatten()

# Create three side-by-side QQ plots
fig, axs = plt.subplots(1, 3, figsize=(18, 5))

# Normalize the data to have zero mean and unit variance
log_returns_daily_normalized = log_returns_daily / np.std(log_returns_daily)

# QQ plot against Student-t distribution with  $\nu = 10$ 
stats.probplot(log_returns_daily_normalized, dist=stats.t, sparams=(10,),
    plot=axs[0])
axs[0].set_title('QQ Plot (Student-t  $\nu=10$ )')
axs[0].set_xlabel('Theoretical Quantiles')
axs[0].set_ylabel('Sample Quantiles')

# QQ plot against Student-t distribution with  $\nu = 5$ 
stats.probplot(log_returns_daily_normalized, dist=stats.t, sparams=(5,),
    plot=axs[1])
axs[1].set_title('QQ Plot (Student-t  $\nu=5$ )')
axs[1].set_xlabel('Theoretical Quantiles')
axs[1].set_ylabel('Sample Quantiles')


# QQ plot against Student-t distribution with  $\nu = 3$ 
stats.probplot(log_returns_daily_normalized, dist=stats.t, sparams=(3,),
    plot=axs[2])
axs[2].set_title('QQ Plot (Student-t  $\nu=3$ )')
axs[2].set_xlabel('Theoretical Quantiles')
axs[2].set_ylabel('Sample Quantiles')

# Adjust spacing between QQ plots
plt.tight_layout()

# Save the plot in png format
plt.savefig('Latex/Img/qqplt_tstudents_AMZNdaily.pdf', format='pdf',
    bbox_inches='tight')

plt.show()

```



Kurtosis of the sample

```
[119]: from scipy.stats import kurtosis
exc_kurt = kurtosis(Rt_d_all) - 3
print("Excess Kurtosis = ", exc_kurt)
```

Excess Kurtosis = 10.51221657872248

The sample of data defined by the simple returns of AMZN stock is **HEAVY TAILED**

12 4.5/High Frequency non-Gaussianity

1. Overall Shapes
2. Lilliefors Test

```
[120]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats as stats
import seaborn as sns

# Extract daily log-returns
log_price_monthly = pt_m_all # Ensure this is a pandas DataFrame or Series
log_returns_monthly = rt_m_all # Ensure this is a pandas DataFrame or Series

# If log_returns_daily is a DataFrame, convert it to a 1D array (assuming
↳ 'column_name' is the name of the column)
log_returns_monthly = log_returns_monthly.values.flatten() # Ensure it's 1D
↳ array

# Calculate monthly log-returns
log_price_yearly = pt_y_all # Ensure this is a pandas DataFrame or Series
log_returns_yearly = rt_y_all # Ensure this is a pandas DataFrame or Series

# If log_returns_monthly is a DataFrame, convert it to a 1D array (assuming
↳ 'column_name' is the name of the column)
log_returns_yearly = log_returns_yearly.values.flatten() # Ensure it's 1D array

# Create the figure with four subplots
fig, axs = plt.subplots(2, 2, figsize=(12, 8))

# Plot histogram of daily log-returns
sns.histplot(log_returns_monthly, bins=30, color='lime', edgecolor='black',
↳ kde_kws={'color': 'red'}, ax=axs[0, 0], stat='density')
```

```

axs[0, 0].plot(np.linspace(log_returns_monthly.min(), log_returns_monthly.
    ↪max(), 100),
                stats.norm.pdf(np.linspace(log_returns_monthly.min(), ↪
    ↪log_returns_monthly.max(), 100),
                                log_returns_monthly.mean(), log_returns_monthly.
    ↪std()), color='red', linewidth=2)
axs[0, 0].set_title('Histogram of Monthly Log>Returns')
axs[0, 0].set_xlabel('Log>Returns')
axs[0, 0].set_ylabel('Density')

# Plot histogram of monthly log-returns
sns.histplot(log_returns_yearly, bins=30, color='lime', edgecolor='black', ↪
    ↪kde_kws={'color': 'red'}, ax=axs[0, 1], stat='density')
axs[0, 1].plot(np.linspace(log_returns_yearly.min(), log_returns_yearly.max(), ↪
    ↪100),
                stats.norm.pdf(np.linspace(log_returns_yearly.min(), ↪
    ↪log_returns_yearly.max(), 100),
                                log_returns_yearly.mean(), log_returns_yearly.
    ↪std()), color='red', linewidth=2)
axs[0, 1].set_title('Histogram of Yearly Log>Returns')
axs[0, 1].set_xlabel('Log>Returns')
axs[0, 1].set_ylabel('Density')

# QQ plot of daily log-returns
stats.probplot(log_returns_yearly, dist="norm", plot=axes[1, 0])
axes[1, 0].set_title('QQ Plot of Monthly Log>Returns')
axes[1, 0].set_xlabel('Normal Quantiles')
axes[1, 0].set_ylabel('Sample Quantiles')

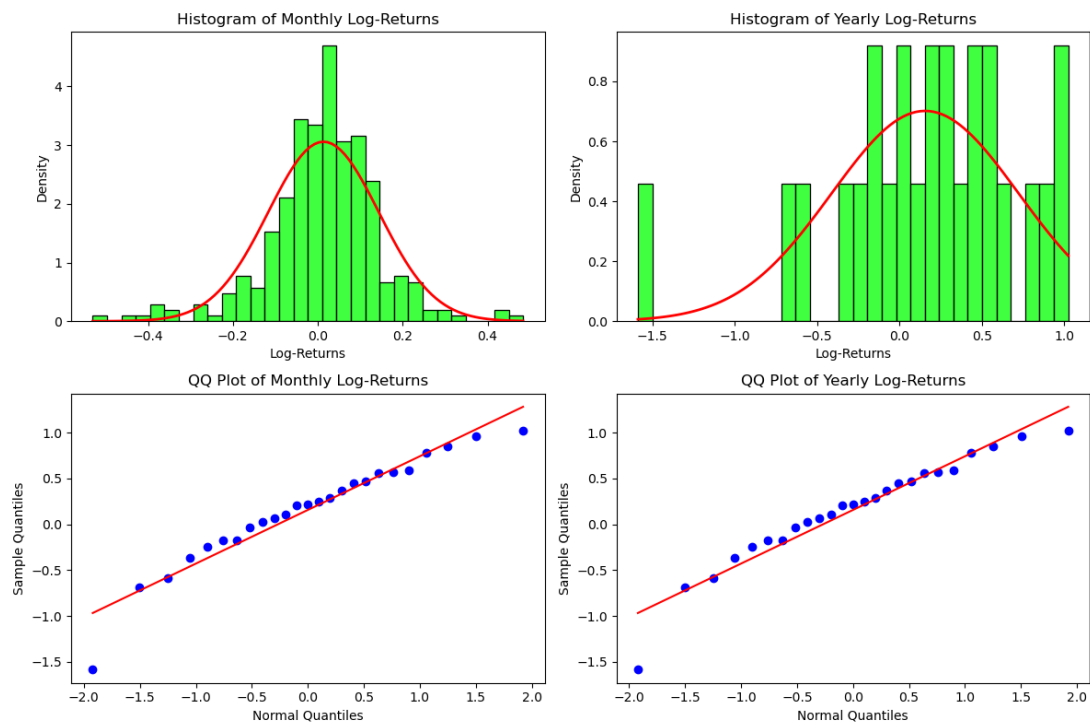
# QQ plot of monthly log-returns
stats.probplot(log_returns_yearly, dist="norm", plot=axes[1, 1])
axes[1, 1].set_title('QQ Plot of Yearly Log>Returns')
axes[1, 1].set_xlabel('Normal Quantiles')
axes[1, 1].set_ylabel('Sample Quantiles')

# Adjust spacing between plots
plt.tight_layout()

# Save the plot in pdf format
plt.savefig('Latex/Img/QQplot_monthly_yearly_AMZN.pdf', format='pdf', ↪
    ↪bbox_inches='tight')

# Show the plot
plt.show()

```



```
[121]: import yfinance as yf
import matplotlib.pyplot as plt
import numpy as np
import scipy.stats as stats
import seaborn as sns

# Compute annual log-returns
log_returns_weekly = rt_w_all

log_returns_yearly = rt_y_all
# Compute mean and std
mean_data = log_returns_yearly.mean()
sd_data = log_returns_yearly.std()
samp_size = len(log_returns_yearly)
seq_ind = np.arange(1, samp_size + 1, 1)
emp_cdf = seq_ind / samp_size
emp_cdf_2 = (seq_ind - 1) / samp_size
my_data_ordered = np.sort(log_returns_yearly)
theor_cdf = stats.norm.cdf(my_data_ordered, mean_data, sd_data)

# Set the layout
fig, axs = plt.subplots(1, 2, figsize=(12, 6))
```



```

# Left panel: empirical and Normal cdf's
sns.ecdfplot(log_returns_yearly, ax=axes[0], label='Empirical CDF')
axes[0].plot(np.linspace(log_returns_yearly.min(), log_returns_yearly.max(),
↳100),
            stats.norm.cdf(np.linspace(log_returns_yearly.min(),
↳log_returns_yearly.max(), 100),
                        mean_data, sd_data),
            color='red', linewidth=2, label='Normal CDF')
axes[0].set_xlabel('AMZN sorted annual log returns')
axes[0].set_ylabel('CDF')
axes[0].set_title('')
axes[0].legend()

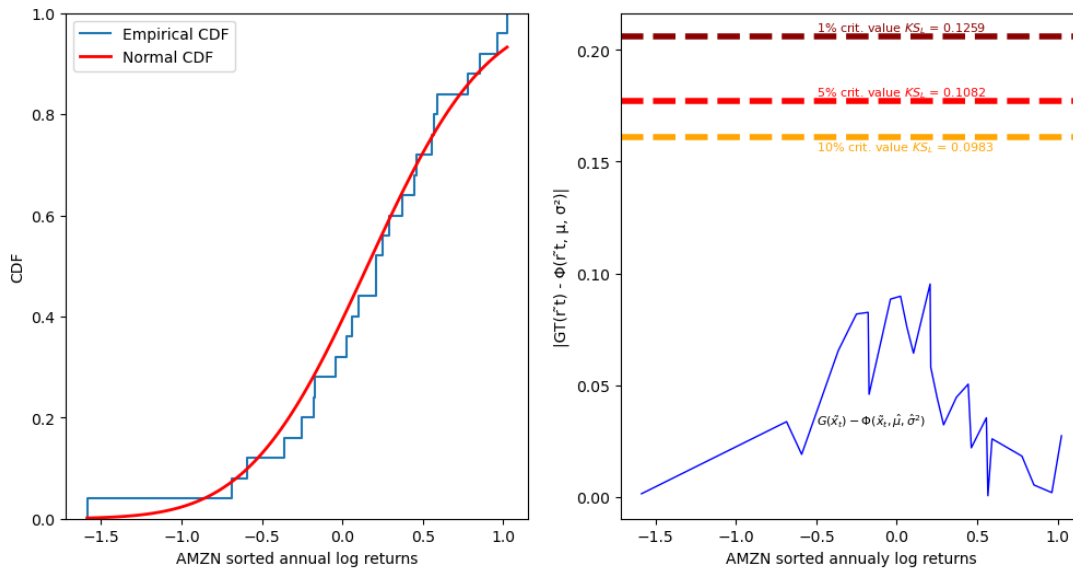
# Right panel: Lilliefors test
KS_L_stat1 = np.max(np.abs(emp_cdf - theor_cdf))
KS_L_stat2 = np.max(np.abs(emp_cdf_2 - theor_cdf))
KS_L_stat = max(KS_L_stat1, KS_L_stat2)
axes[1].plot(my_data_ordered, np.abs(emp_cdf_2 - theor_cdf), color='blue',
↳linewidth=1)
axes[1].axhline(y=0.805/np.sqrt(samp_size), color='orange', linewidth=4,
↳linestyle='--')
axes[1].axhline(y=0.886/np.sqrt(samp_size), color='red', linewidth=4,
↳linestyle='--')
axes[1].axhline(y=1.031/np.sqrt(samp_size), color='darkred', linewidth=4,
↳linestyle='--')
axes[1].text(-0.5, 0.805/np.sqrt(samp_size)-0.006, '10% crit. value $KS_L$ = 0.
↳0983', fontsize=8, color='orange')
axes[1].text(-0.5, 0.886/np.sqrt(samp_size)+0.002, '5% crit. value $KS_L$ = 0.
↳1082', fontsize=8, color='red')
axes[1].text(-0.5, 1.031/np.sqrt(samp_size)+0.002, '1% crit. value $KS_L$ = 0.
↳1259', fontsize=8, color='darkred')
axes[1].text(-0.5, 0.032, '$G(\tilde{x}_t) - \Phi(\tilde{x}_t, \hat{\mu},
↳\hat{\sigma}^2)$', fontsize=8)
axes[1].set_xlabel('AMZN sorted annual log returns')
axes[1].set_ylabel('$|G(\tilde{r}_t) - \Phi(\tilde{r}_t, \hat{\mu}, \hat{\sigma}^2)|$')
axes[1].set_title('')

# Set the space within plots
plt.tight_layout()

# Save the figure in png format
plt.savefig('Latex/Img/lillie_test_AMZNannually.pdf', format='pdf',
↳bbox_inches='tight')

plt.show()

```



13 4.6/ Returns are not autocorrelated

1. Daily, Weekly and Monthly Autocorrelations

```
[122]: import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import acf

# Calculate empirical autocorrelations for daily, weekly, monthly, and yearly
↳ returns
lags = 40

# Daily ACF
acf_daily_values = acf(Rt_d_all, nlags=lags)

# Weekly ACF
acf_weekly_values = acf(Rt_w_all, nlags=lags)

# Monthly ACF
acf_monthly_values = acf(Rt_m_all, nlags=lags)

# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(Rt_d_all))

# Create the autocorrelation plot with Bartlett intervals for each time frame
plt.figure(figsize=(12, 8))
```

```

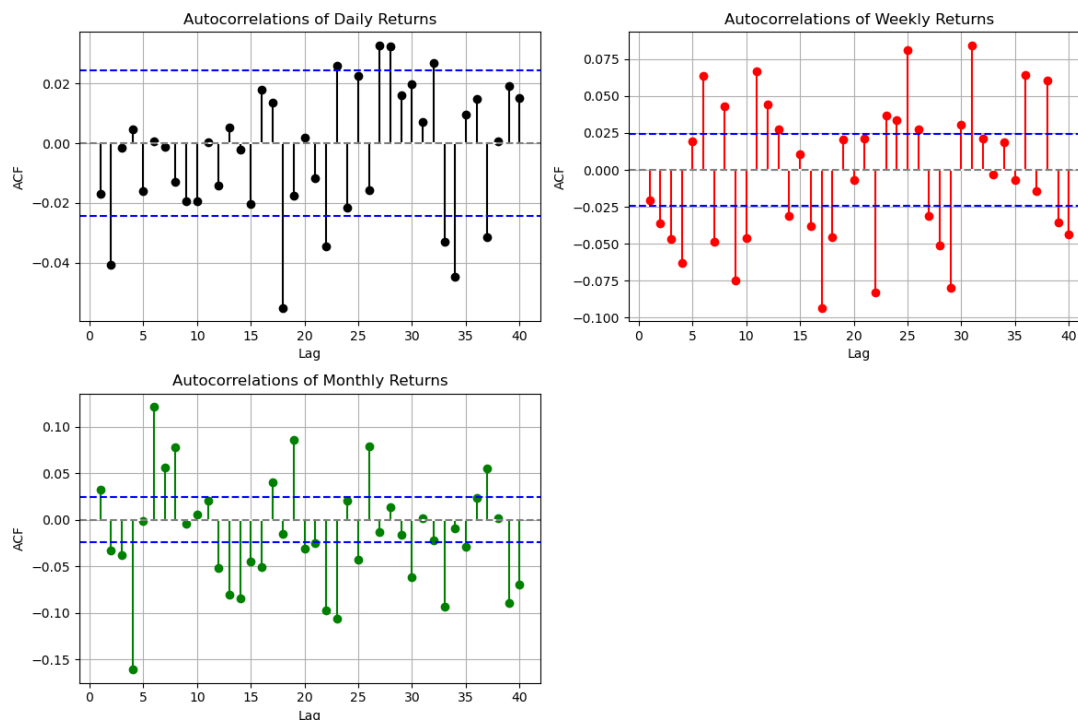
# Plot daily autocorrelations
plt.subplot(2, 2, 1)
plt.stem(np.arange(1, lags + 1), acf_daily_values[1:], linefmt='k-',
        ↪markerfmt='ko', basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
plt.title('Autocorrelations of Daily Returns')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)

# Plot weekly autocorrelations
plt.subplot(2, 2, 2)
plt.stem(np.arange(1, lags + 1), acf_weekly_values[1:], linefmt='r-',
        ↪markerfmt='ro', basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
plt.title('Autocorrelations of Weekly Returns')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)

# Plot monthly autocorrelations
plt.subplot(2, 2, 3)
plt.stem(np.arange(1, lags + 1), acf_monthly_values[1:], linefmt='g-',
        ↪markerfmt='go', basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
plt.title('Autocorrelations of Monthly Returns')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)

# Adjust layout and show plot
plt.tight_layout()
plt.show()

```



```
[123]: """autocorrelate = pt_d_all.shift(1).corrwith(pt_d_all, method='pearson')
print(autocorrelate.round(4))"""
```

```
[123]: "autocorrelate = pt_d_all.shift(1).corrwith(pt_d_all,
method='pearson')\nprint(autocorrelate.round(4))"
```

14 4.7/ Returns feature volatility clustering long run range dependence of squared returns

```
[124]: # Extract daily log-returns
log_returns_daily = rt_d_all

# Parameter for the empirical autocorrelation
lags = 40

# Creation of the three side-by-side graphs
fig, axs = plt.subplots(1, 4, figsize=(30, 6))

# ACF of daily log-returns with confidence bands
acf_values_daily = acf(abs(log_returns_daily), nlags=lags)
confint = 1.96 / np.sqrt(len(log_returns_daily))
confint_upper = np.full(lags, confint)
```

```

confint_lower = -np.full(lags, confint)

axs[0].stem(np.arange(1, lags + 1), acf_values_daily[1:], linefmt='k-',
            ↪markerfmt='ko', basefmt='w-')
axs[0].axhline(y=0, color='gray', linestyle='--')
axs[0].plot(np.arange(1, lags + 1), confint_upper, color='blue',
            ↪linestyle='dashed')
axs[0].plot(np.arange(1, lags + 1), confint_lower, color='blue',
            ↪linestyle='dashed')
axs[0].set_ylim(-0.1, 0.3)
axs[0].set_title('ACF - Daily Absolute Log>Returns')
axs[0].set_xlabel('Lag')
axs[0].set_ylabel('ACF')
axs[0].grid(True)

# ACF of weekly log-returns with confidence bands
acf_values_weekly = acf(abs(log_returns_weekly), nlags=lags)
confint_weekly = 1.96 / np.sqrt(len(log_returns_weekly))
confint_weekly_upper = np.full(lags, confint_weekly)
confint_weekly_lower = -np.full(lags, confint_weekly)

axs[1].stem(np.arange(1, lags + 1), acf_values_weekly[1:], linefmt='k-',
            ↪markerfmt='ko', basefmt='w-')
axs[1].axhline(y=0, color='gray', linestyle='--')
axs[1].plot(np.arange(1, lags + 1), confint_weekly_upper, color='blue',
            ↪linestyle='dashed')
axs[1].plot(np.arange(1, lags + 1), confint_weekly_lower, color='blue',
            ↪linestyle='dashed')
axs[1].set_ylim(-0.1, 0.3)
axs[1].set_title('ACF - Weekly Absolute Log>Returns')
axs[1].set_xlabel('Lag')
axs[1].set_ylabel('ACF')
axs[1].grid(True)

# ACF of monthly log-returns with confidence bands
acf_values_monthly = acf(abs(log_returns_monthly), nlags=lags)
confint_monthly = 1.96 / np.sqrt(len(log_returns_monthly))
confint_monthly_upper = np.full(lags, confint_monthly)
confint_monthly_lower = -np.full(lags, confint_monthly)

axs[2].stem(np.arange(1, lags + 1), acf_values_monthly[1:], linefmt='k-',
            ↪markerfmt='ko', basefmt='w-')
axs[2].axhline(y=0, color='gray', linestyle='--')
axs[2].plot(np.arange(1, lags + 1), confint_monthly_upper, color='blue',
            ↪linestyle='dashed')

```

```

axs[2].plot(np.arange(1, lags + 1), confint_monthly_lower, color='blue',
            linestyle='dashed')
axs[2].set_ylim(-0.13, 0.39)
axs[2].set_title('ACF - Monthly Absolute Log>Returns')
axs[2].set_xlabel('Lag')
axs[2].set_ylabel('ACF')
axs[2].grid(True)

# ACF of annual log-returns with confidence bands
lags = 24

acf_values_yearly = acf(abs(log_returns_yearly), nlags=lags)
confint = 1.96 / np.sqrt(len(log_returns_yearly))
confint_upper = np.full(lags, confint)
confint_lower = -np.full(lags, confint)

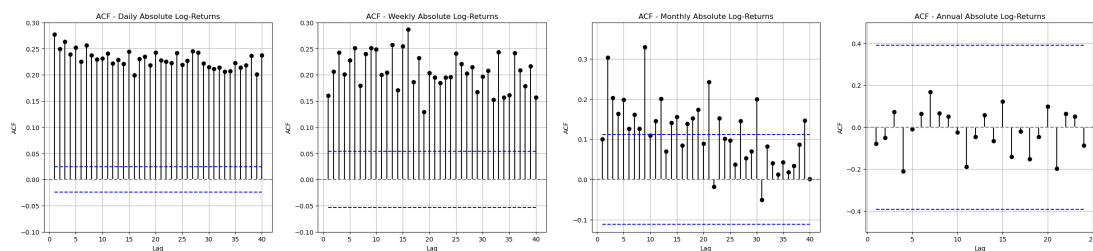
axs[3].stem(np.arange(1, lags + 1), acf_values_yearly[1:], linefmt='k-',
            markerfmt='ko', basefmt='w-')
axs[3].axhline(y=0, color='gray', linestyle='--')
axs[3].plot(np.arange(1, lags + 1), confint_upper, color='blue',
            linestyle='dashed')
axs[3].plot(np.arange(1, lags + 1), confint_lower, color='blue',
            linestyle='dashed')
axs[3].set_ylim(-0.5, 0.5)
axs[3].set_title('ACF - Annual Absolute Log>Returns')
axs[3].set_xlabel('Lag')
axs[3].set_ylabel('ACF')
axs[3].grid(True)

# Adjusting the spacing between graphs
plt.tight_layout()

# Save the graphic in png format
plt.savefig('Latex/Img/Fact7_AbsoluteLogReturns.pdf', format='pdf',
            bbox_inches='tight')

plt.show()

```



15 4.8/ Leverage Effect

```
[ ]: # Define a function
def ccf(x, y, lag_max = 100):
    # Compute correlation
    result = ss.correlate(y - np.mean(y), x - np.mean(x), method='direct') / (
        np.std(y) * np.std(x) * len(y))
    # Define the length
    length = (len(result) - 1) // 2
    lo = length - lag_max
    hi = length + (lag_max + 1)
    return result[lo:hi]

# Choose the max lag and execute the function
lag_max = 10
log_returns_daily = np.array(log_returns_daily)
cross_corr = ccf(log_returns_daily, log_returns_daily**2, lag_max=lag_max)

# Plot results
lags = np.arange(-lag_max, lag_max + 1)

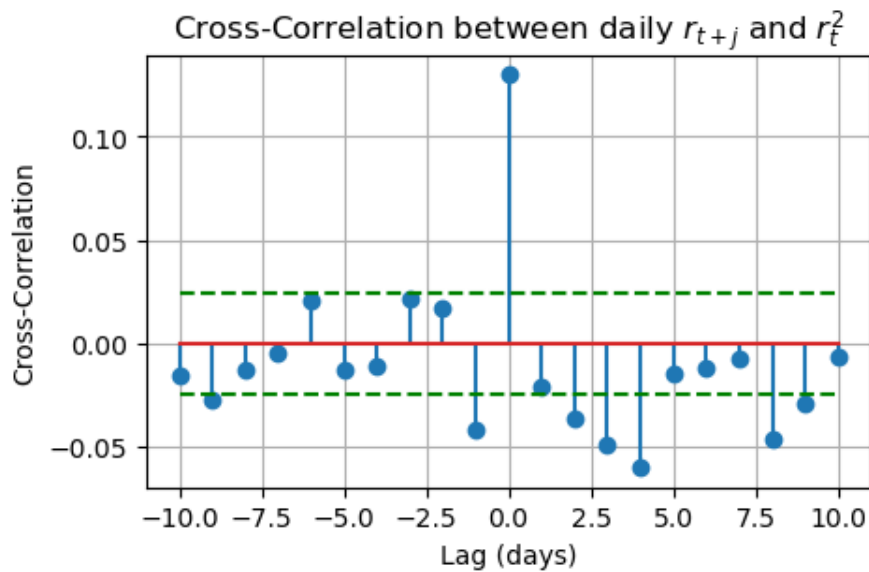
# ACF of monthly log-returns with confidence bands
confint_daily = 1.96 / np.sqrt(len(log_returns_daily))
confint_daily_upper = np.full(len(lags), confint_daily)
confint_daily_lower = -np.full(len(lags), confint_daily)

plt.figure(figsize=(5, 3))
plt.stem(lags, cross_corr)
plt.plot(lags, confint_daily_upper, color='green', linestyle='dashed')
plt.plot(lags, confint_daily_lower, color='green', linestyle='dashed')
plt.xlabel('Lag (days)')
plt.ylabel('Cross-Correlation')
plt.title('Cross-Correlation between daily  $r_{t+j}$  and  $r_t^2$ ')
plt.grid(True)

# Add the bartlett intervals

plt.savefig('Latex/Img/Fact8_CrossCorr_r_r2.pdf', format='pdf',
            bbox_inches='tight')

plt.show()
```



```
[126]: print(Pt_d_all, type)
```

```
Date
1999-01-21      2.650000
1999-01-22      3.075000
1999-01-25      2.809375
1999-01-26      2.877344
1999-01-27      3.140625
...
2024-10-09     185.169998
2024-10-10     186.649994
2024-10-11     188.820007
2024-10-14     187.539993
2024-10-15     187.690002
Name: Pt.d, Length: 6476, dtype: float64 <class 'type'>
```

```
[127]: #Get the starting and ending date of our stock
start_date = Pt_d_all.index.min()
end_date = Pt_d_all.index.max()

# Get VIX data
VIX = yf.download("^VIX", start=start_date, end=end_date)

# Extract and Rename the adjusted closing prices
VIX_d = VIX["Adj Close"]
VIX_d.name = 'VIX.d'
```



```

# Mutate the Index into a DatetimeIndex
VIX_d.index = pd.to_datetime(VIX_d.index)

# Merge the two datasets and rename columns
merged_df = pd.merge(Pt_d_all, VIX_d, on='Date', how='outer') # outer: only
↳ common indexes (dates)
merged_df.head()

# Compute changes in pt and VIX compared to previous period (NaN are kept)
diff_df = merged_df.diff()
diff_df.head()

# Remove from the price dataframe
merged_df = merged_df.dropna()
# And from the second one
diff_df = diff_df.dropna()

# Define the figure parameters
fig, ax1 = plt.subplots(figsize=(10, 3))

# Customizing x-axis labels for December of each year
date_labels = pd.date_range(start=start_date, end=end_date, freq='3Y')
formatted_labels = [f'Dec-{date.year}' for date in date_labels]
# Add label and rotate them
plt.xticks(date_labels, formatted_labels, rotation=45)

# Work on the first y-axis: S&P
ax1.plot(merged_df.index, merged_df['Pt.d'], label="AMZN" + ' Prices',
↳ color='blue')
ax1.set_xlabel('Date')
ax1.set_ylabel("AMZN", color='blue')
ax1.tick_params(axis='y', labelcolor='blue')

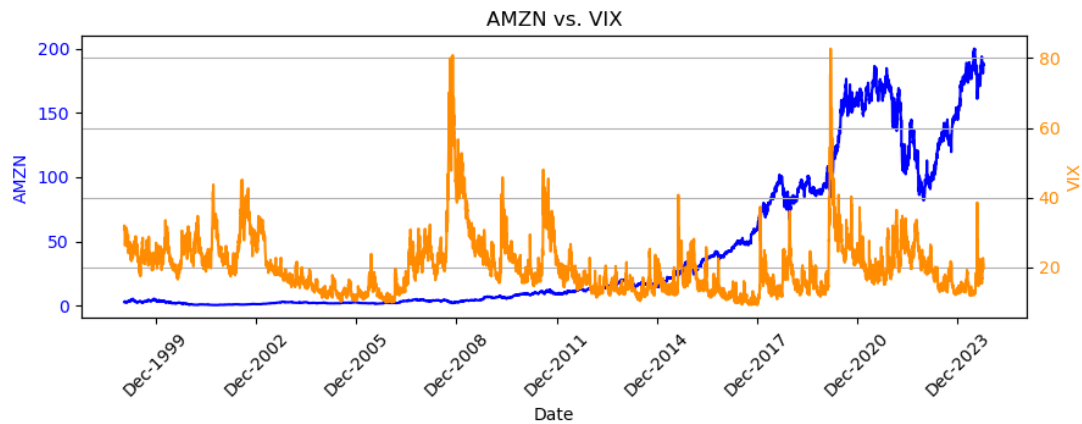
# Work on the second y-axis: VIX
ax2 = ax1.twinx()
ax2.plot(merged_df.index, merged_df['VIX.d'], label='VIX', color='darkorange')
ax2.set_ylabel('VIX', color='darkorange')
ax2.tick_params(axis='y', labelcolor='darkorange')

# Adjust the figure
plt.title("AMZN" + ' vs. VIX')
plt.grid(True)

# Save the figure
plt.savefig('Latex/Img/Fact8.pdf', format='png', bbox_inches='tight')
plt.show()

```

[*****100%*****] 1 of 1 completed



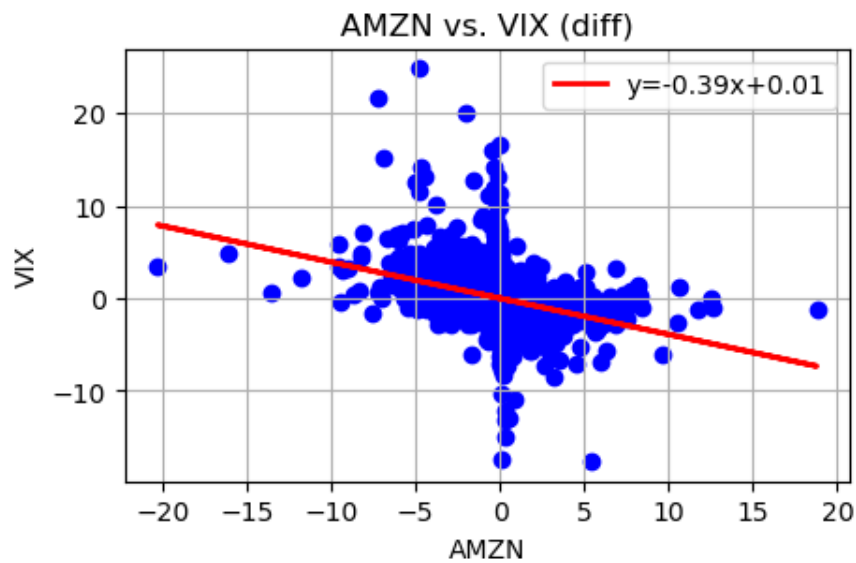
```
[128]: plt.figure(figsize=(5, 3))
plt.scatter(diff_df['Pt.d'], diff_df['VIX.d'], color='blue', marker='o')

# Add labels and title
plt.xlabel("AMZN")
plt.ylabel('VIX')
plt.title("AMZN" + ' vs. VIX (diff)')
plt.grid(True)

# Add regression line
coefficients = np.polyfit(diff_df['Pt.d'], diff_df['VIX.d'], 1)
regression_line = np.polyval(coefficients, diff_df['Pt.d'])
plt.plot(diff_df['Pt.d'], regression_line, color='red', linewidth=2,
         label='y='+str(round(coefficients[0],2))+ 'x'+str(round(coefficients[1],2)))
plt.legend()

plt.savefig('Latex/Img/Fact_8_3'+ "AMZN" + '_.pdf', format='pdf',
         bbox_inches='tight')

# Show plot
plt.show()
```



15.0.1 Other Material

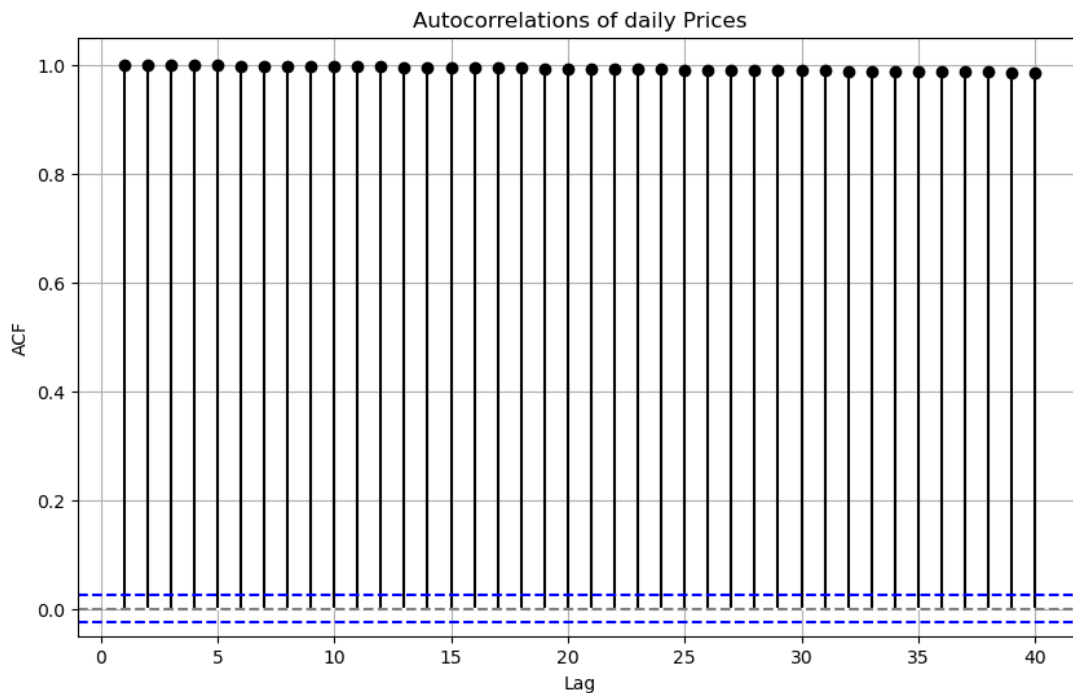
took from stylized facts 1

```
[129]: from statsmodels.tsa.stattools import acf

# Calculate empirical autocorrelation
lags = 40
acf_values = acf(pt_d_all, nlags=lags)

# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(pt_d_all))

# Create the autocorrelation plot with Bartlett intervals
plt.figure(figsize=(10, 6))
plt.stem(np.arange(1, lags + 1), acf_values[1:], linefmt='k-', markerfmt='ko',
         basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
plt.title('Autocorrelations of daily Prices')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)
# plt.savefig('Latex/Autocorrel_daily.pdf', format='pdf', bbox_inches='tight')
plt.show()
```



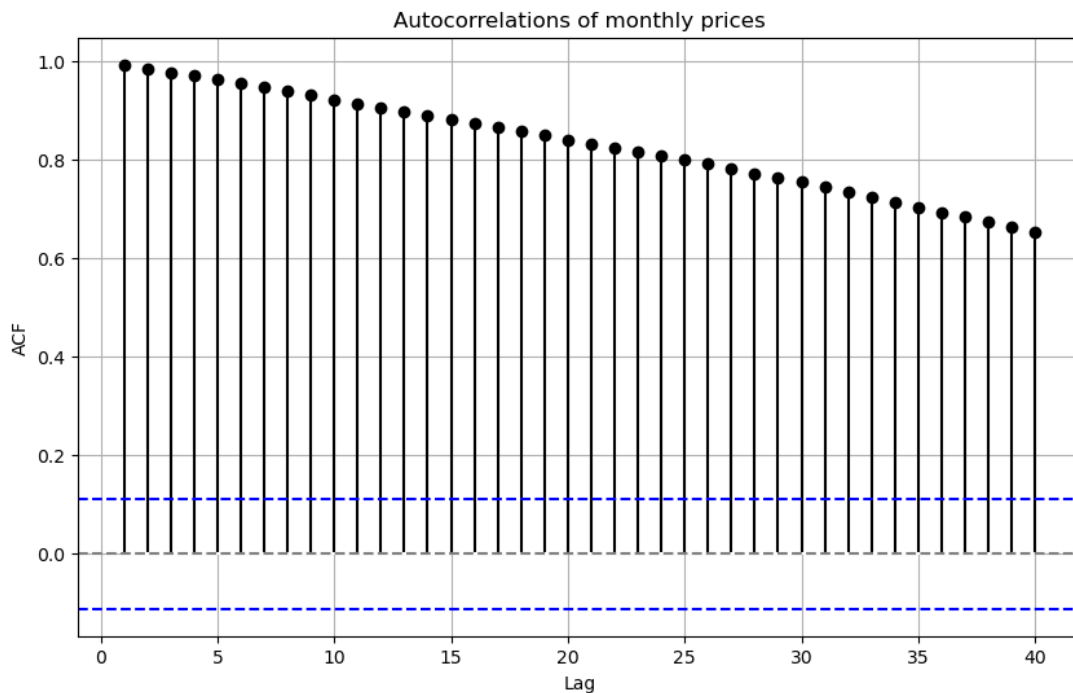
ACF with monthly data

```
[130]: from statsmodels.tsa.stattools import acf

# Calculate empirical autocorrelation
lags = 40
acf_values = acf(pt_m_all, nlags=lags)

# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(pt_m_all))

# Create the autocorrelation plot with Bartlett intervals
plt.figure(figsize=(10, 6))
plt.stem(np.arange(1, lags + 1), acf_values[1:], linefmt='k-', markerfmt='ko',
        basefmt='w-')
plt.axhline(y=0, color='gray', linestyle='--')
plt.axhline(y=Bart_Int, color='blue', linestyle='--')
plt.axhline(y=-Bart_Int, color='blue', linestyle='--')
plt.title('Autocorrelations of monthly prices')
plt.xlabel('Lag')
plt.ylabel('ACF')
plt.grid(True)
# plt.savefig('Latex/Autocorrel_monthly.pdf', format='pdf', bbox_inches='tight')
plt.show()
```



15.0.2 The two figures subplotted

```
[131]: fig, axs = plt.subplots(1, 2, figsize=(18, 6))
#first fig
# Calculate empirical autocorrelation
lags = 40
acf_values = acf(pt_d_all, nlags=lags)

# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(pt_d_all))

axs[0].stem(np.arange(1, lags + 1), acf_values[1:], linefmt='k-',
            markerfmt='ko', basefmt='w-')
axs[0].axhline(y=0, color='gray', linestyle='--')
axs[0].axhline(y=Bart_Int, color='blue', linestyle='--')
axs[0].axhline(y=-Bart_Int, color='blue', linestyle='--')
axs[0].set_title('Autocorrelations of daily Prices')
axs[0].set_xlabel('Lag')
axs[0].set_ylabel('ACF')
axs[0].grid(True)
#second fig
# Calculate empirical autocorrelation
lags = 40
acf_values = acf(pt_m_all, nlags=lags)
```

```

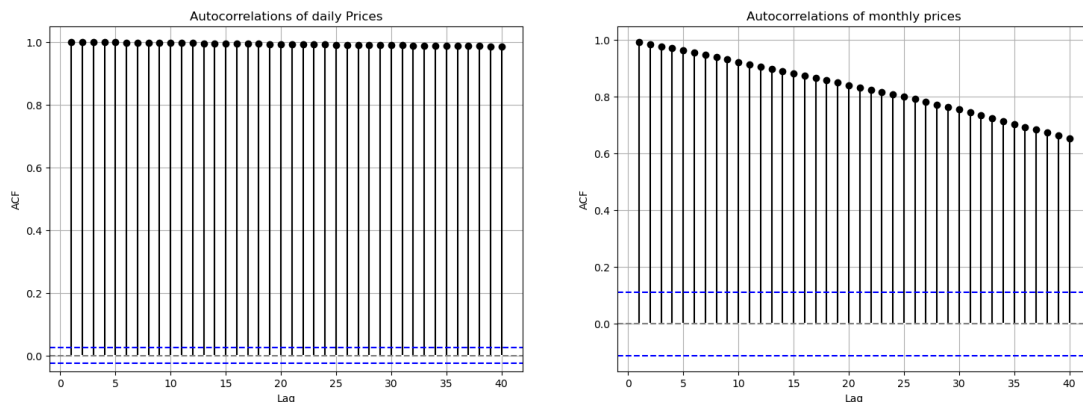
# Calculate Bartlett intervals
Bart_Int = 1.96 / np.sqrt(len(pt_m_all))

axs[1].stem(np.arange(1, lags + 1), acf_values[1:], linefmt='k-',
            ↪markerfmt='ko', basefmt='w-')
axs[1].axhline(y=0, color='gray', linestyle='--')
axs[1].axhline(y=Bart_Int, color='blue', linestyle='--')
axs[1].axhline(y=-Bart_Int, color='blue', linestyle='--')
axs[1].set_title('Autocorrelations of monthly prices')
axs[1].set_xlabel('Lag')
axs[1].set_ylabel('ACF')
axs[1].grid(True)

plt.savefig('Latex/Img/Autocorrel_daily_monthly.pdf', format='pdf',
            ↪bbox_inches='tight')

plt.show()

```



Histogram of daily prices and normal density

```

[132]: # Set up the subplots
fig, axs = plt.subplots(1, 2, figsize=(18, 7))

# Histogram and Normal Distribution (Daily)
axs[0].hist(rt_d_all, bins=50, density=True, color="lightgreen")
norm_y = stats.norm.pdf(np.linspace(rt_d_all.min(), rt_d_all.max()), loc=np.
    ↪mean(rt_d_all), scale=np.std(rt_d_all))
axs[0].plot(np.linspace(rt_d_all.min(), rt_d_all.max()), norm_y, color="blue",
            ↪linewidth=1)
axs[0].set_xlabel("daily log-return")
axs[0].set_title("Histogram and Normal Distribution (Daily)")

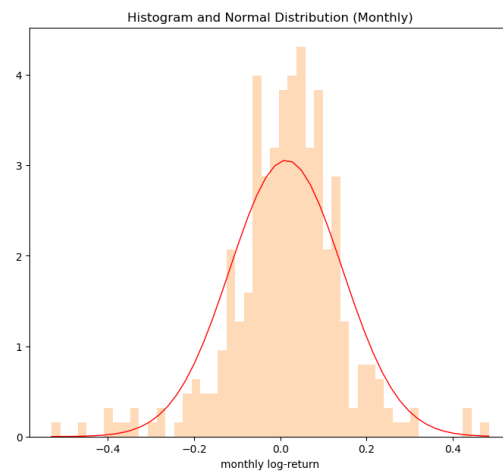
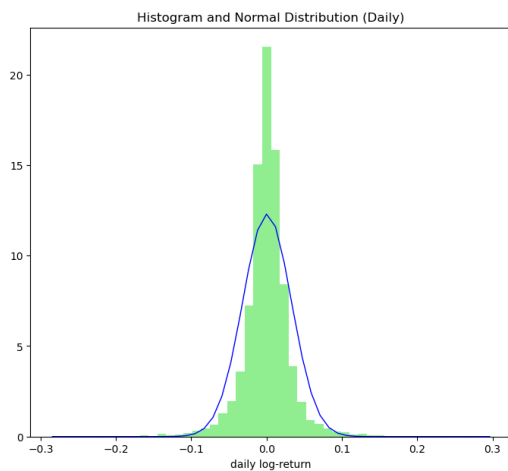
```

```

# Histogram and Normal Distribution (Monthly)
axs[1].hist(rt_m_all, bins=50, density=True, color="peachpuff")
norm_y = stats.norm.pdf(np.linspace(rt_m_all.min(), rt_m_all.max()), loc=np.
    ↪mean(rt_m_all), scale=np.std(rt_m_all))
axs[1].plot(np.linspace(rt_m_all.min(), rt_m_all.max()), norm_y, color="red",
    ↪linewidth=1)
axs[1].set_xlabel("monthly log-return")
axs[1].set_title("Histogram and Normal Distribution (Monthly)")

# Adjust layout and display the plot
#plt.tight_layout()
plt.savefig('Latex/Img/Histogram_and_normal_distrib_daily.pdf', format='pdf',
    ↪bbox_inches='tight')
plt.show()

```



15.0.3 QQ-plot (Normal distribution)

```

[133]: # Set up the subplots
fig, axs = plt.subplots(1, 2, figsize=(12, 5))

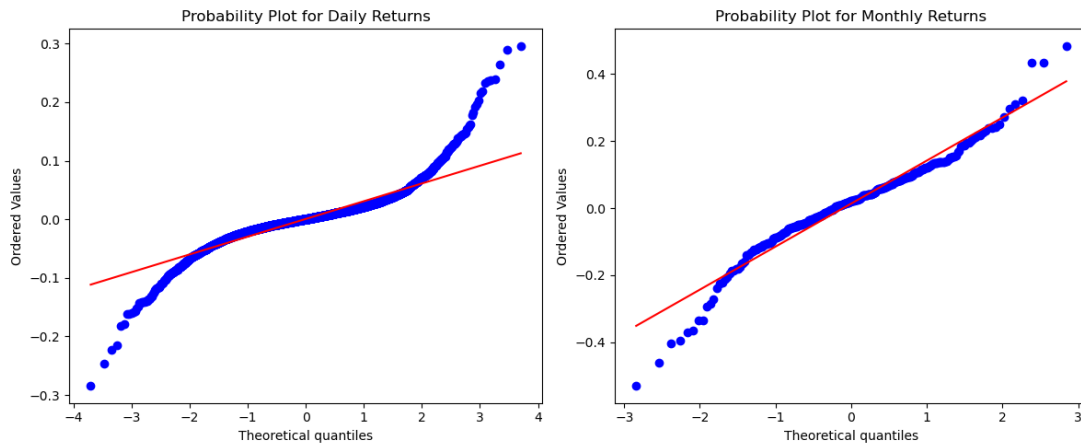
# Probability Plot for Daily Returns
stats.probplot(rt_d_all, dist="norm", plot=axs[0])
axs[0].set_title("Probability Plot for Daily Returns")

# Probability Plot for Monthly Returns
stats.probplot(rt_m_all, dist="norm", plot=axs[1])
axs[1].set_title("Probability Plot for Monthly Returns")

# Adjust layout and display the plot

```

```
plt.tight_layout()
#plt.savefig('Latex/QQ-plot.pdf', format='pdf', bbox_inches='tight')
plt.show()
```



Check how the QQ-plots change aggregating data

```
[134]: # Set up the subplots
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Probability Plot for Daily Returns
stats.probplot(rt_d_all, dist="norm", plot=axes[0, 0])
axes[0, 0].set_title("Probability Plot for Daily Returns")
axes[0, 0].grid(True)

# Probability Plot for Weekly Returns
stats.probplot(rt_w_all, dist="norm", plot=axes[0, 1])
axes[0, 1].set_title("Probability Plot for Weekly Returns")
axes[0, 1].grid(True)

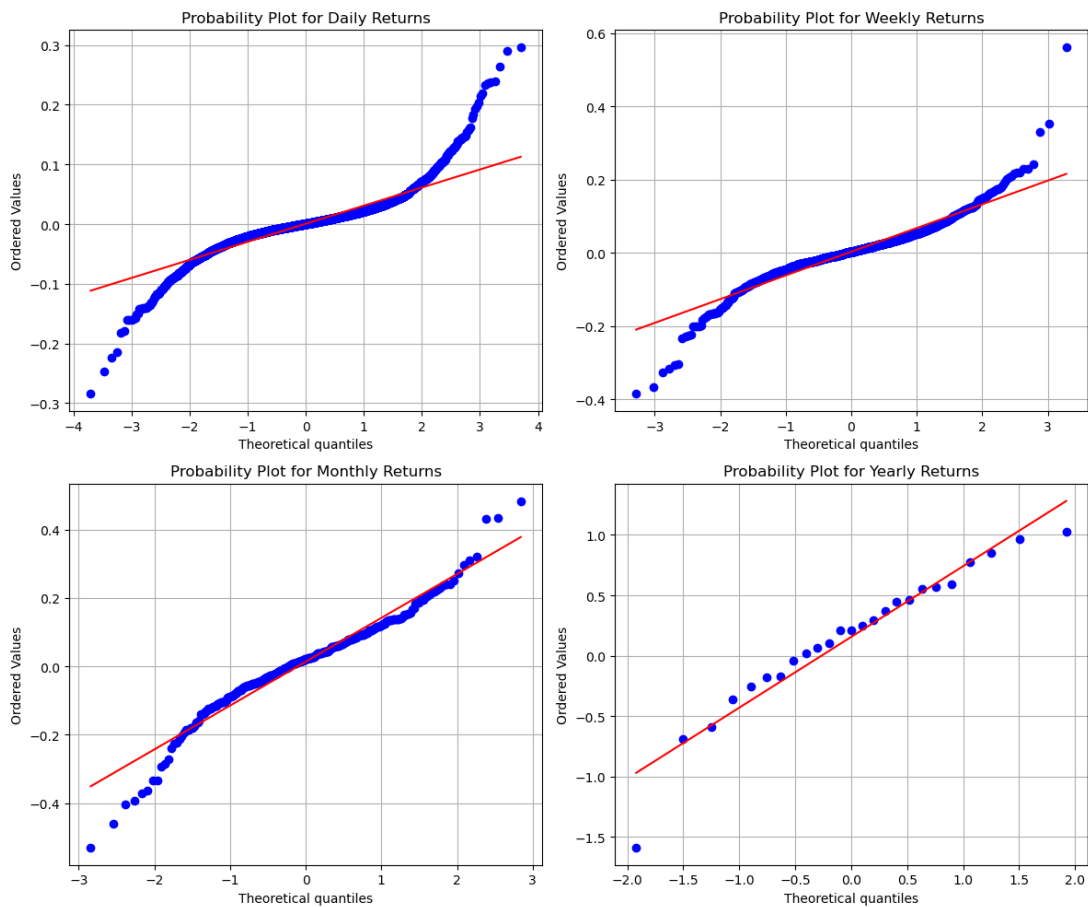
# Probability Plot for Monthly Returns
stats.probplot(rt_m_all, dist="norm", plot=axes[1, 0])
axes[1, 0].set_title("Probability Plot for Monthly Returns")
axes[1, 0].grid(True)

# Probability Plot for Yearly Returns
stats.probplot(rt_y_all, dist="norm", plot=axes[1, 1])
axes[1, 1].set_title("Probability Plot for Yearly Returns")
axes[1, 1].grid(True)

# Adjust layout and display the plot
plt.tight_layout()
```



```
plt.show()
```



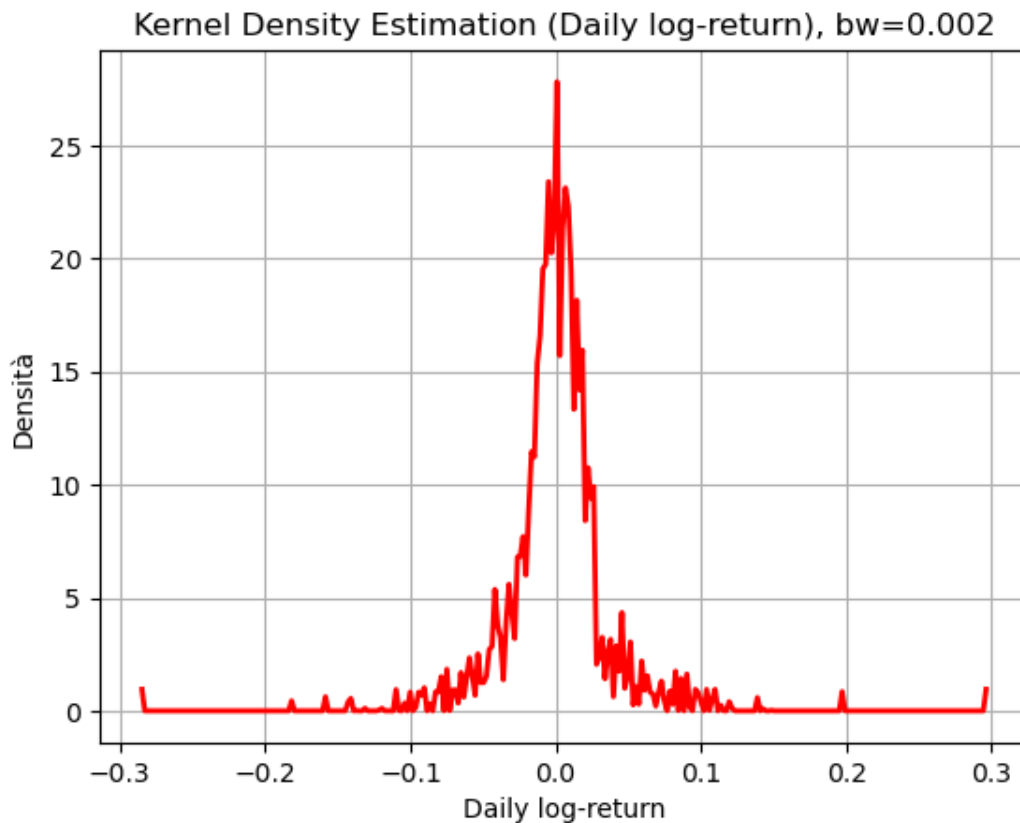
15.0.4 Kernel density

It is similar to a smooth histogram!

```
[135]: ## Compute the kernel density: daily returns
# divide the interval between the min and max returns into 300 segments
density_eval_points = np.linspace(rt_d_all.min(), rt_d_all.max(), num=300)
# estimate the kernel density of our returns
kde = gaussian_kde(rt_d_all, bw_method=0.002)
# and evaluate in the interval defined above
density_estimation = kde(density_eval_points)

# Plotting
plt.plot(density_eval_points, density_estimation, color='red', lw=2,
        label='Kernel density')
plt.xlabel("Daily log-return")
```

```
plt.ylabel("Densità")
plt.title("Kernel Density Estimation (Daily log-return), bw=0.002")
plt.grid(True)
plt.show()
```



bw_method defines the *bandwidth parameter*:

⇒ the larger the bandwidth, the smoother the histogram:

```
[136]: ## Compute the kernel density: daily returns
# divide the interval between the min and max returns into 300 segments
density_eval_points = np.linspace(rt_d_all.min(), rt_d_all.max(), num=300)
# estimate the kernel density of our returns
kde = gaussian_kde(rt_d_all, bw_method=0.05)
# and evaluate in the interval defined above
density_estimation = kde(density_eval_points)

# Empirical mean and std
mean_empirical= log_returns_daily.mean()
std_empirical= log_returns_daily.std()
```

```

x=np.arange(-0.3,0.3,0.01)

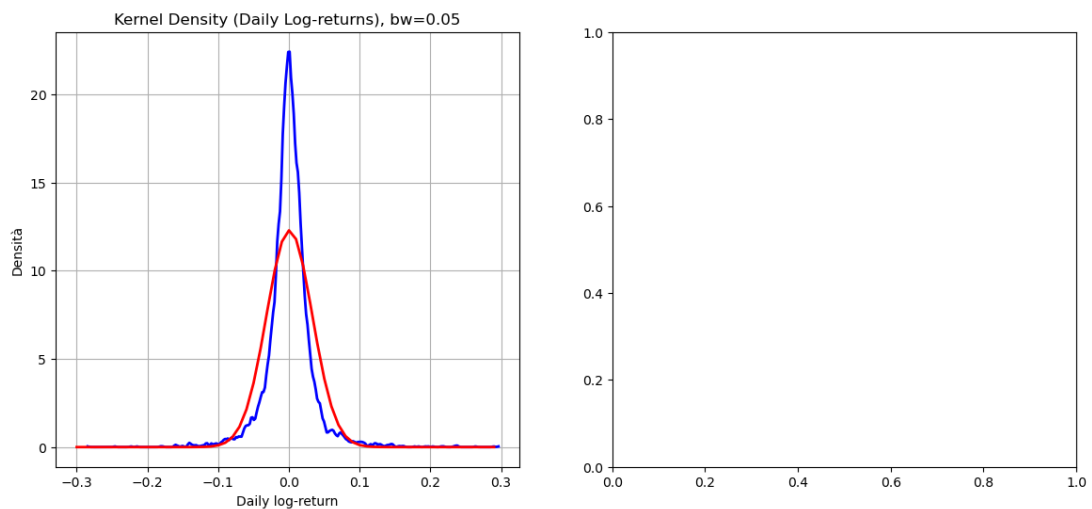
fig,axs=plt.subplots(1,2,figsize=(14,6))
# Plotting

#sns.kdeplot(log_returns_daily, color='blue', ax=axs[0])

# 1rst plot is kernel density daily log returns, compared to the standard
↪normal
axs[0].plot(density_eval_points, density_estimation, color='blue', lw=2,
↪label='Kernel density')
axs[0].plot(x, stats.norm.pdf(x, mean_empirical, std_empirical), color='red',
↪linewidth=2)
axs[0].set_xlabel("Daily log-return")
axs[0].set_ylabel("Densità")
axs[0].set_title("Kernel Density (Daily Log-returns), bw=0.05")
axs[0].grid(True)

"""sns.histplot(log_returns_daily, bins=60, color='lime', edgecolor='black',
↪kde_kws={'color': 'red'}, stat='density', ax=axs[1])
axs[1].plot(stats.norm.pdf(np.linspace(log_returns_daily.min(),
↪log_returns_daily.max(), 100),log_returns_daily.mean(), log_returns_daily.
↪std()),color='red', linewidth=2)
"""
plt.show()

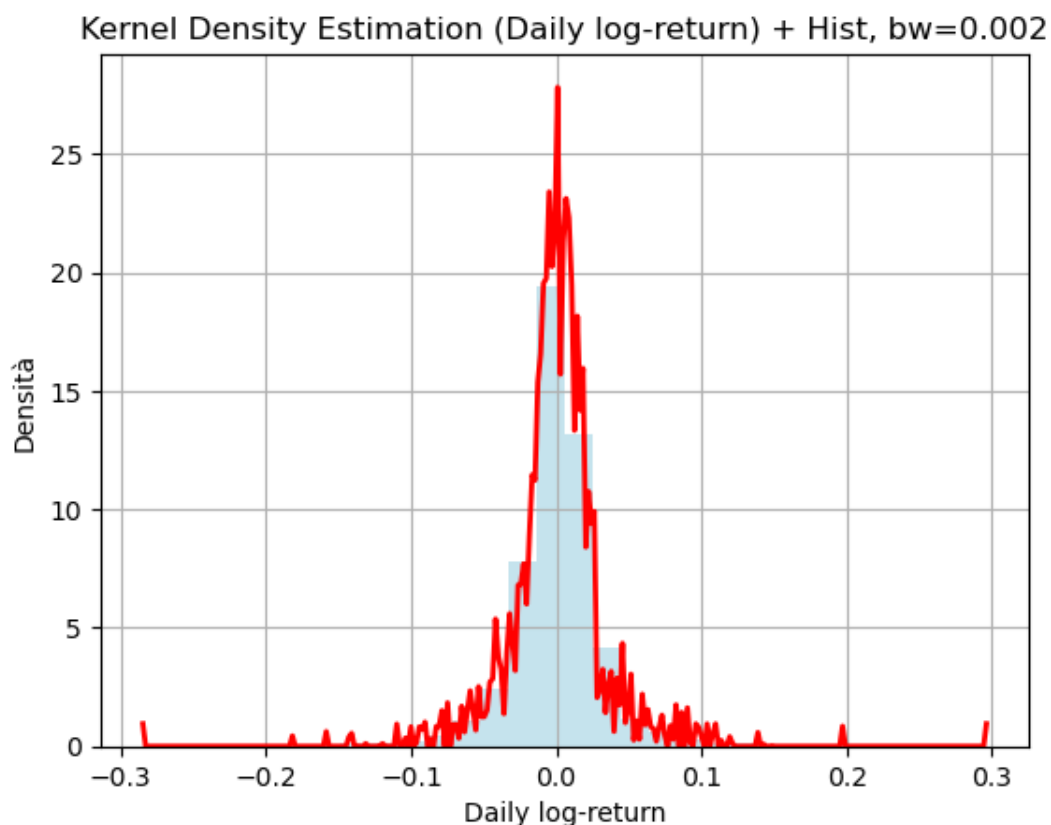
```



Above, there was an issue with a histogram, but we did not use it in the interpretation


```
[137]: ## Compute the kernel density: daily returns
# divide the interval between the min and max returns into 300 segments
density_eval_points = np.linspace(rt_d_all.min(), rt_d_all.max(), num=300)
# estimate the kernel density of our returns
kde = gaussian_kde(rt_d_all, bw_method=0.002)
# and evaluate in the interval defined above
density_estimation = kde(density_eval_points)

# Plotting
plt.hist(rt_d_all, bins=30, density=True, alpha=0.7, color='lightblue')
plt.plot(density_eval_points, density_estimation, color='red', lw=2,
        label='Kernel density')
plt.xlabel("Daily log-return")
plt.ylabel("Densità")
plt.title("Kernel Density Estimation (Daily log-return) + Hist, bw=0.002")
plt.grid(True)
plt.show()
```



```
[138]: ## Compute the kernel density: daily returns
# divide the interval between the min and max returns into 300 segments
density_eval_points = np.linspace(rt_d_all.min(), rt_d_all.max(), num=300)
# estimate the kernel density of our returns
kde = gaussian_kde(rt_d_all, bw_method=0.002)
# and evaluate in the interval defined above
density_estimation = kde(density_eval_points)
# on the same interval, we evaluate a Normal pdf
pdf_theoretical = norm.pdf(density_eval_points, np.mean(rt_d_all), np.
    ↪std(rt_d_all))

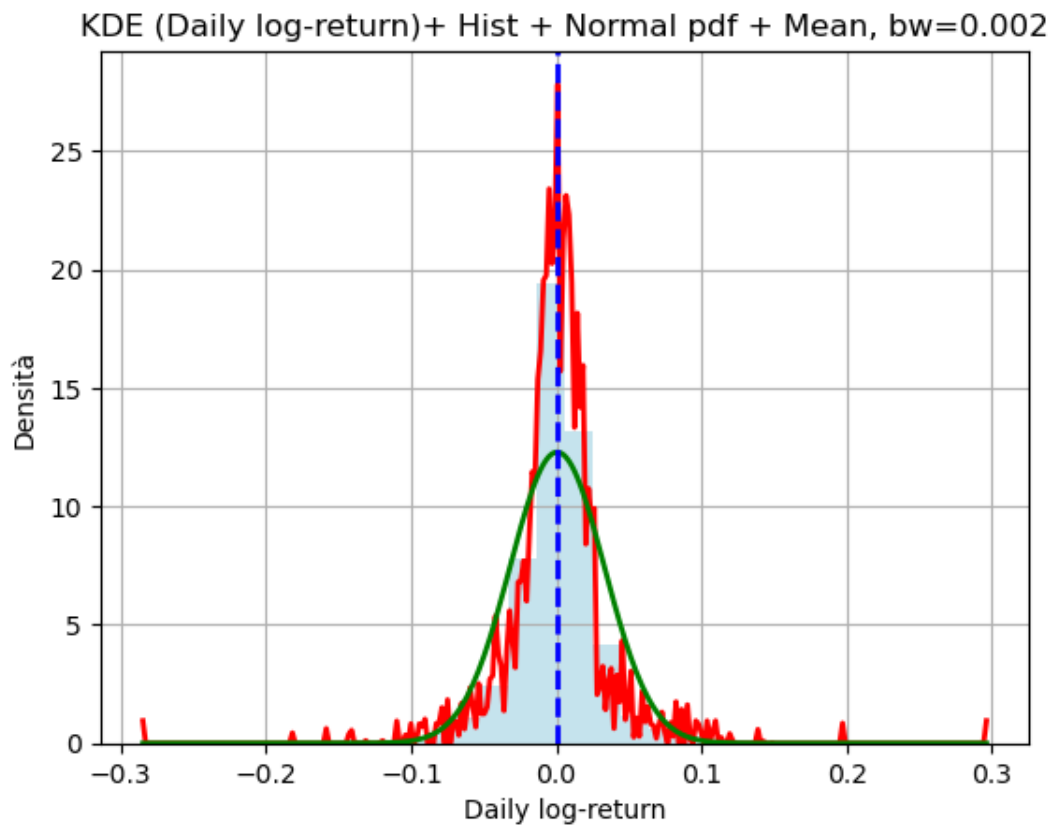
# Plotting
plt.hist(rt_d_all, bins=30, density=True, alpha=0.7, color='lightblue')
plt.plot(density_eval_points, density_estimation, color='red', lw=2,
    ↪label='Kernel density')
plt.plot(density_eval_points, pdf_theoretical, color='green', lw=2, label='PDF
    ↪Teorica (Normale)')
plt.xlabel("Daily log-return")
plt.ylabel("Densità")
plt.title("KDE (Daily log-return)+ Hist + Normal pdf, bw=0.002")
plt.grid(True)
plt.show()
```



```
[139]: ## Compute the kernel density: daily returns
# divide the interval between the min and max returns into 300 segments
density_eval_points = np.linspace(rt_d_all.min(), rt_d_all.max(), num=300)
# estimate the kernel density of our returns
kde = gaussian_kde(rt_d_all, bw_method=0.002)
# and evaluate in the interval defined above
density_estimation = kde(density_eval_points)
# on the same interval, we evaluate a Normal pdf
pdf_theoretical = norm.pdf(density_eval_points, np.mean(rt_d_all), np.
    ↪std(rt_d_all))
# compute the mean
mean_data = np.mean(rt_d_all)

# Plotting
plt.hist(rt_d_all, bins=30, density=True, alpha=0.7, color='lightblue')
plt.plot(density_eval_points, density_estimation, color='red', lw=2,
    ↪label='Kernel density')
plt.plot(density_eval_points, pdf_theoretical, color='green', lw=2, label='PDF
    ↪Teorica (Normale)')
plt.axvline(mean_data, color='blue', linestyle='dashed', linewidth=2,
    ↪label='Media')

plt.xlabel("Daily log-return")
plt.ylabel("Densità")
plt.title("KDE (Daily log-return)+ Hist + Normal pdf + Mean, bw=0.002")
plt.grid(True)
plt.show()
```



15.1 Summary Statistics

There is a function which compute some summary statistics...not really the ones we want called describe:

```
[140]: rt_d_all.describe()
```

```
[140]: count      6475.000000
      mean         0.000658
      std          0.032465
      min         -0.284568
      25%         -0.012598
      50%          0.000413
      75%          0.013969
      max          0.296181
      Name: rt_d_all, dtype: float64
```

15.1.1 Skewness & Kurtosis

We use the fucntions which came from scipy.stats:

from scipy.stats import gaussian_kde, norm, iqr, skew, kurtosis, jarque_bera, kstest, anderson

These functions replicate the formulas you find on slides.

```
[141]: rt_d_skew = skew(rt_d_all, nan_policy='omit')
rt_d_kurt = kurtosis(rt_d_all, nan_policy='omit')

print("The skewness is:", rt_d_skew)
print("The kurtosis is:", rt_d_kurt)

# NOTE: There are several formulas to compute skewness and kurtosis.
#       These functions both divide the summations of the estimators by 1/T
```

The skewness is: 0.4304836875303971

The kurtosis is: 11.128625063290052

Aggregational Kurtosis We compute the kurtosis of the daily, weekly, monthly, and annual returns:

```
[142]: rt_d_kurt = kurtosis(rt_d_all, nan_policy='omit')
rt_w_kurt = kurtosis(rt_w_all, nan_policy='omit')
rt_m_kurt = kurtosis(rt_m_all, nan_policy='omit')
rt_y_kurt = kurtosis(rt_y_all, nan_policy='omit')

print("Daily: ", round(rt_d_kurt,3))
print("Weekly: ", round(rt_w_kurt,3))
print("Monthly: ", round(rt_m_kurt,3))
print("Annual: ", round(rt_y_kurt,3))
```

Daily: 11.129

Weekly: 7.605

Monthly: 2.604

Annual: 1.461

15.1.2 Normality Tests

Compute Normality tests and sample summary statistics

Jarque-Bera Test

```
[143]: JB_rt_d = jarque_bera(rt_d_all)
# first position (0): statistic
print("JB Stat: ", round(JB_rt_d[0],3))
# second position (1): p-value
print("JB p-value: ", JB_rt_d[1])
```

JB Stat: 33612.686

JB p-value: 0.0

Check the Aggregational Normality:


```
[144]: print("JB p-value", "daily", "returns:", jarque_bera(rt_d_all)[1])
print("JB p-value", "weekly", "returns:", jarque_bera(rt_w_all)[1])
print("JB p-value", "monthly", "returns:", jarque_bera(rt_m_all)[1])
print("JB p-value", "yearly", "returns:", jarque_bera(rt_y_all)[1])
```

```
JB p-value daily returns: 0.0
JB p-value weekly returns: 0.0
JB p-value monthly returns: 0.0
JB p-value yearly returns: 0.04262486815078237
```

We can also compute the p-value. The JB Stats follows a χ^2_2 distribution. So:

```
[145]: p_value = 1 - stats.chi2.cdf(STATISTIC, df=2)
print("The associated p-value is:", p_value)
```

```
The associated p-value is: 0.04262486815078237
```

15.1.3 Other normality tests:

Lilliefors test:

```
[146]: lill_rt_y = lilliefors(rt_y_all)
print("Stat:", lill_rt_y[0])
print("p-val:", lill_rt_y[1])
```

```
Stat: 0.0952201438460884
p-val: 0.7974329823750796
```

Kolmogorov-Smirnov test:

```
[147]: ks_rt_y = kstest(rt_y_all, 'norm')
print("Stat:", ks_rt_y[0])
print("p-val:", ks_rt_y[1])
```

```
Stat: 0.24100976414208733
p-val: 0.0919870853397472
```

Anderson-Darling test:

```
[148]: ad_rt_y = anderson(rt_y_all, 'norm')
print("Stat:", ad_rt_y[0])
print("critical val:", ad_rt_y[1])
print("sign level:", ad_rt_y[2])
```

```
Stat: 0.34577306424633036
critical val: [0.514 0.586 0.703 0.82 0.975]
sign level: [15. 10. 5. 2.5 1.]
```

15.2 Generates table exactly equal to the one in slide n.91

Personalized table of summary statistics.

```
[149]: # X contains returns at different frequencies
```

```
X = {  
    'daily': rt_d_all,  
    'weekly': rt_w_all,  
    'monthly': rt_m_all,  
    'annual': rt_y_all  
}
```

```
[150]: def multi_fun(x):
```

```
    stat_tab = {  
        'Mean': round(np.mean(x) * 100,5),  
        'St.Deviation': round(np.std(x) * 100,5),  
        'Diameter.C.I.Mean': round(1.96 * np.sqrt(np.var(x) / len(x)) * 100,5),  
        'Skewness': round(skew(x),5),  
        'Kurtosis': round(kurtosis(x),5),  
        'Excess.Kurtosis': round(kurtosis(x) - 3,5),  
        'Min': round(np.min(x) * 100,5),  
        'Quant5': round(np.quantile(x, 0.05) * 100,5),  
        'Quant25': round(np.quantile(x, 0.25) * 100,5),  
        'Median': round(np.quantile(x, 0.50) * 100,5),  
        'Quant75': round(np.quantile(x, 0.75) * 100,5),  
        'Quant95': round(np.quantile(x, 0.95) * 100,5),  
        'Max': round(np.max(x) * 100,5),  
        'Jarque.Bera.stat': round(jarque_bera(x)[0],5),  
        'Jarque.Bera.pvalue.X100': round(jarque_bera(x)[1] * 100,5),  
        'Lillie.test.stat': round(lilliefors(x)[0],5),  
        'Lillie.test.pvalue.X100': round(lilliefors(x)[1] * 100,5),  
        'N.obs': len(x)  
    }  
    return stat_tab
```

1. Define a new dictionary to store the stats:
 - a. key will contains the key (i.e., daily, weekly, ...)
 - b. data will contains the returns
2. Apply *multi_fun* to each data series
3. Define a DataFrame with the stats results
4. Print the dictionary

```
[151]: # 1.
```

```
statistics_dict = {}
```

```
# 2.
```

```
statistics_dict = {  
    key: multi_fun(data.iloc[1:])  
}
```

```

    for key, data in X.items()
}
# apply multi_fun to each returns ("series" in pandas)
# which is located in one of the four key of our dictionary X
# 3.
statistics_df = pd.DataFrame(statistics_dict)

# 4.
print(statistics_df)

```

	daily	weekly	monthly	annual
Mean	0.06351	0.31014	1.32164	22.85692
St.Deviation	3.24131	6.76830	13.06275	45.28052
Diameter.C.I.Mean	0.07896	0.36213	1.45887	18.11598
Skewness	0.41992	0.05008	-0.45895	-0.15137
Kurtosis	11.15158	7.60655	2.59462	-0.64793
Excess.Kurtosis	8.15158	4.60655	-0.40538	-3.64793
Min	-28.45678	-38.51804	-53.02674	-68.54809
Quant5	-4.61051	-9.74288	-20.16713	-55.72274
Quant25	-1.25994	-2.64062	-4.98163	-7.23999
Median	0.04108	0.30519	2.09626	23.07665
Quant75	1.39659	3.40897	8.45973	55.96192
Quant95	4.47118	10.67416	20.90661	94.77653
Max	29.61811	56.11507	48.35221	102.44636
Jarque.Bera.stat	33735.75720	3235.87866	97.20696	0.51147
Jarque.Bera.pvalue.X100	0.00000	0.00000	0.00000	77.43487
Lillie.test.stat	0.10194	0.09591	0.08194	0.06494
Lillie.test.pvalue.X100	0.10000	0.10000	0.10000	99.00000
N.obs	6474.00000	1342.00000	308.00000	24.00000

Export it as a latex table

```

[152]: latex_table = statistics_df.to_latex(index=True)
with open("Latex/8stylized.tex", "w") as file:
    file.write(latex_table)

```

```

/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/805359341.py:1:
FutureWarning: In future versions `DataFrame.to_latex` is expected to utilise
the base implementation of `Styler.to_latex` for formatting and rendering. The
arguments signature may therefore change. It is recommended instead to use
`DataFrame.style.to_latex` which also contains additional functionality.
    latex_table = statistics_df.to_latex(index=True)

```

```

[153]: #skewness & kurtosis dict

def skewness_dict(x):
    stat_tab = {
        'Skewness': round(skew(x),5),
        'Kurtosis': round(kurtosis(x),5),
    }

```

```

}
return stat_tab

```

```

[154]: # Y contains returns at different frequencies
Y = {
    'daily': Rt_d_all,
    'weekly': Rt_w_all,
    'monthly': Rt_m_all,
    'annual': Rt_y_all
}

```

Creating a table with isolated skewness and kurtosis

```

[155]: statistics_dict_sk = {}

statistics_dict_sk_log = {
    key: skewness_dict(data.iloc[1:])
    for key, data in X.items()
}

statistics_dict_sk_simple = {
    key: skewness_dict(data.iloc[1:])
    for key, data in Y.items()
}

#printing results

print("Log returns",pd.DataFrame(statistics_dict_sk_log))
print("Simple returns",pd.DataFrame(statistics_dict_sk_simple))

```

Log returns		daily	weekly	monthly	annual
Skewness	0.41992	0.05008	-0.45895	-0.15137	
Kurtosis	11.15158	7.60655	2.59462	-0.64793	

Simple returns		daily	weekly	monthly	annual
Skewness	1.07310	1.16787	0.40997	0.69139	
Kurtosis	13.54753	13.70696	2.93127	-0.30325	

Export it as a Latex table

```

[156]: latex_table = pd.DataFrame(statistics_dict_sk_log).to_latex(index=True)
with open("Latex/table_Skewness_kurtosis.tex", "w") as file:
    file.write(latex_table)

```

```

/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/3564160307.py:1
: FutureWarning: In future versions `DataFrame.to_latex` is expected to utilise
the base implementation of `Styler.to_latex` for formatting and rendering. The
arguments signature may therefore change. It is recommended instead to use
`DataFrame.style.to_latex` which also contains additional functionality.

```

```
latex_table = pd.DataFrame(statistics_dict_sk_log).to_latex(index=True)
```

```
[157]: # Compute Box Pierce and Ljung Box tests
```

```

for log_returns in [log_returns_daily, log_returns_monthly]:
    my_max_lag = 25
    lags_all = np.arange(1, my_max_lag + 1)
    my_acf = sm.tsa.acf(log_returns, nlags=my_max_lag)
    my_acf_diameter = 1.96 / np.sqrt(len(log_returns))
    my_acf_tstat_0 = (my_acf[1:] - 0) / np.sqrt(1 / len(log_returns))
    my_LjungBox = sm.stats.diagnostic.acorr_ljungbox(log_returns, lags=lags_all,
↪boxpierce=False)
    my_BoxPierce = sm.stats.diagnostic.acorr_ljungbox(log_returns, lags=lags_all,
↪boxpierce=True)
    crit_value_5_BP = stats.chi2.ppf(0.95, lags_all)

my_table = np.column_stack((
    lags_all,
    my_acf[1:],
    np.full(my_max_lag, my_acf_diameter),
    my_acf_tstat_0,
    my_BoxPierce['bp_stat'],
    my_BoxPierce['bp_pvalue'],
    my_LjungBox['lb_stat'],
    my_LjungBox['lb_pvalue'],
    np.full(my_max_lag, crit_value_5_BP)
))

column_names = ["lag", "acf", "acf diam.", "acf test", "B-P stat", "B-P pval",
↪ "L-B stat", "L-B pval", "crit"]

my_table_df = pd.DataFrame(data=my_table, columns=column_names)

# Reducing the selection of lags
my_table_df = my_table_df.iloc[[0,4,14,24]]

# Print the rounded table
my_table_df = my_table_df.round(3)
my_table_df['lag'] = my_table_df['lag'].astype(int)
print(my_table_df)

```

[illegible]

```
4    11.070
14   24.996
24   37.652
```

Export it as a Latex Table

```
[158]: latex_table = my_table_df.to_latex(index=True)
       with open("Latex/LB_BP.tex", "w") as file:
           file.write(latex_table)
```

```
/var/folders/5r/ft807c7n1ngd3fpt2_gwsg0m0000gn/T/ipykernel_78356/4088390180.py:1
: FutureWarning: In future versions `DataFrame.to_latex` is expected to utilise
the base implementation of `Styler.to_latex` for formatting and rendering. The
arguments signature may therefore change. It is recommended instead to use
`DataFrame.style.to_latex` which also contains additional functionality.
  latex_table = my_table_df.to_latex(index=True)
```